

TEST

PYNVME QPAIRS

0000:3d:00.0 Q00:

0000:3d:00.0 Q01: >>> CAZ-82256-Q1 NVMe LITEON 256GB

PYTHON

- scripts
- conformance
- 01_admin_cmd
- 02_nvme_cmd
- 03_features
- 04_registers
- 05_controller
- arbitration_test.py
- interrupt_test.py
- prp_test.py
- sq_cq_test.py
- sqe_cqe_test.py
- 06_tcg
- ftl
- performance
- stress
- dirty_power_cycle_test.py
- endurance_test.py
- test_ioworker_jedec_e...
- test_ioworker_jedec_e...
- C test_replay_jedec_clien...
- marathon_test.py
- trace
- test_examples.py

CMDLOG 0000:3d:00.0 Q00 X

1 2020-06-22 08:24:29.897073 [cmd001: Identify]
 2 0x007d0006, 0x00000001, 0x00000000, 0x00000000
 3 0x00000000, 0x00000000, 0x4ac0f000, 0x00000001
 4 0x00000000, 0x00000000, 0x00000000, 0x00000000
 5 0x00000000, 0x00000000, 0x00000000, 0x00000000
 6 2020-06-22 08:24:29.897102: [cpl: SUCCESS]
 7 0x00000000, 0x00000000, 0x0000000c, 0x0001007d
 8
 9 2020-06-22 08:24:29.896836 [cmd002: Identify]
 10 0x007d0006, 0x00000001, 0x00000000, 0x00000000
 11 0x00000000, 0x00000000, 0x4ac0f000, 0x00000001
 12 0x00000000, 0x00000000, 0x00000000, 0x00000000
 13 0x00000000, 0x00000000, 0x00000000, 0x00000000
 14 2020-06-22 08:24:29.897059: [cpl: SUCCESS]
 15 0x00000000, 0x00000000, 0x0000000b, 0x0001007d
 16
 17 2020-06-22 08:24:28.452258 [cmd003: Format NVM]
 18 0x007d0080, 0x00000001, 0x00000000, 0x00000000
 19 0x00000000, 0x00000000, 0x00000000, 0x00000000
 20 0x00000000, 0x00000000, 0x00000000, 0x00000000
 21 0x00000000, 0x00000000, 0x00000000, 0x00000000
 22 2020-06-22 08:24:29.573384: [cpl: SUCCESS]
 23 0x00000000, 0x00000000, 0x0000000a, 0x0001007d
 24

Performance Gauge X

...
 pynvme MB/s 104.96
 IOPS K/s 009.72

PROBLEMS OUTPUT ... Python Test Log

root@... /home/craneclay/pynvme, python -m pytest -v

plugins: cov-2.10.0

collected 1 item

scripts/stress/endurance_test.
 py::test_replay_jedec_client_trace
 ----- live log setup

 [2020-06-22 08:24:27.666] INFO script(60): setup random seed:
 0x4400ee0c
 ----- live log call

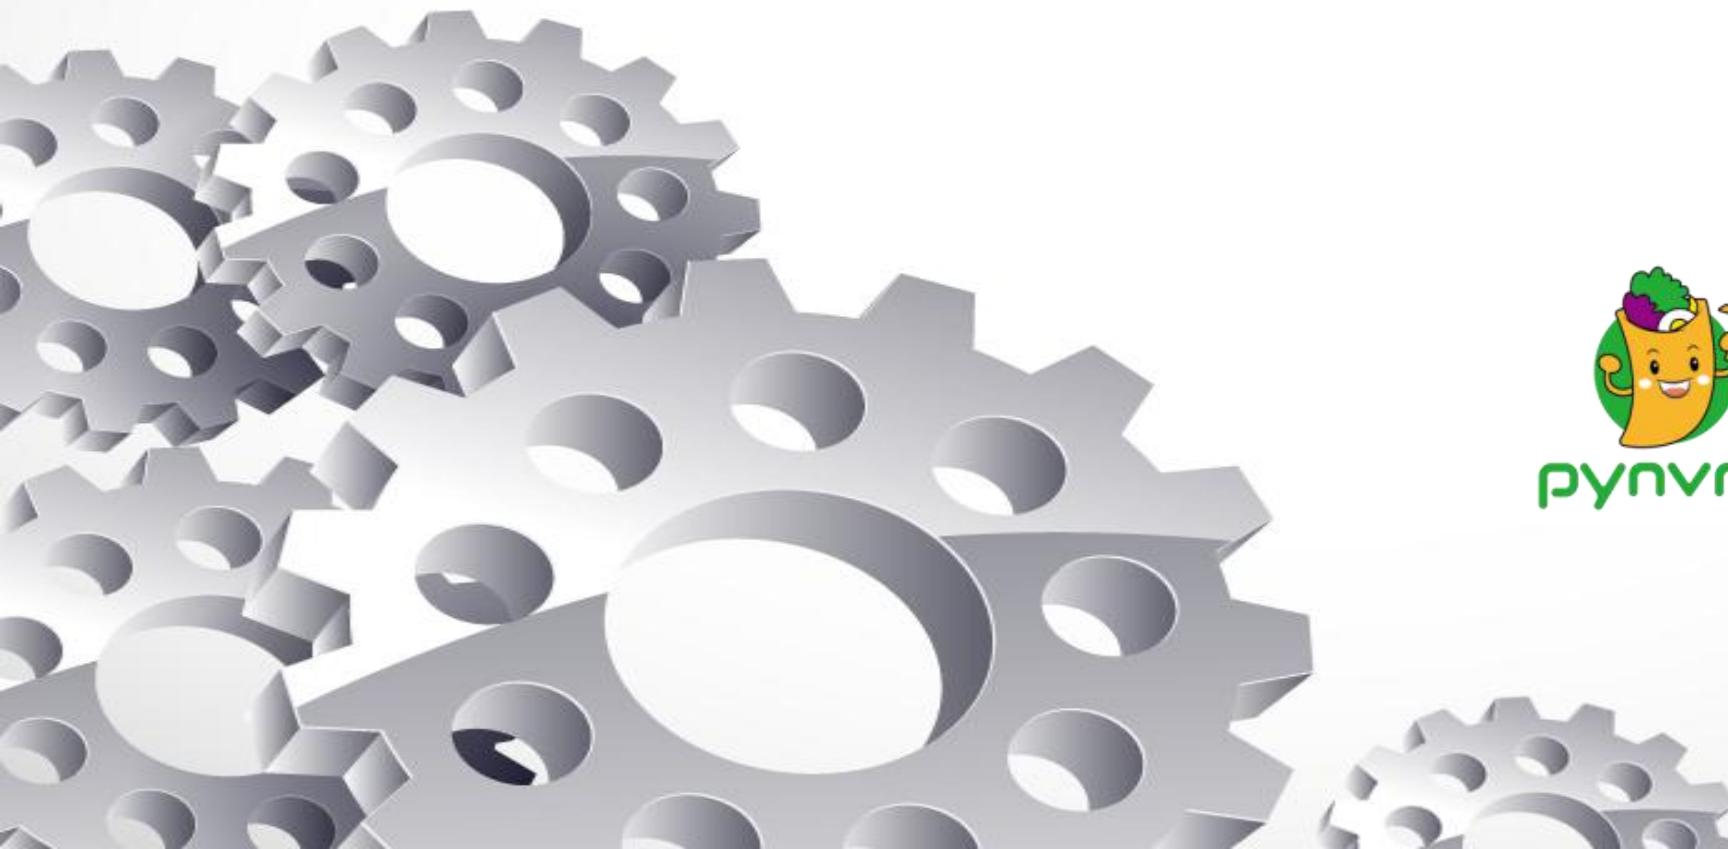
 [2020-06-22 08:24:32.194] INFO test_replay_jedec_client_trace
 (95): replay batch 0
 [2020-06-22 08:24:41.628] INFO test_replay_jedec_client_trace
 (95): replay batch 1
 [2020-06-22 08:24:45.792] INFO test_replay_jedec_client_trace
 (95): replay batch 2
 [2020-06-22 08:25:12.507] INFO test_replay_jedec_client_trace
 (95): replay batch 3
 [2020-06-22 08:25:31.836] INFO test_replay_jedec_client_trace
 (95): replay batch 4
 [2020-06-22 08:25:41.889] INFO test_replay_jedec_client_trace
 (95): replay batch 5
 [2020-06-22 08:25:45.449] INFO test_replay_jedec_client_trace
 (95): replay batch 6
 [2020-06-22 08:25:55.715] INFO test_replay_jedec_client_trace
 (95): replay batch 7
 [2020-06-22 08:26:09.297] INFO test_replay_jedec_client_trace



pynvme builds your own tests.

master* Python 3.8.3 64-bit 0 △ 0 Running Tests / Ln 1, Col 1 Spaces: 4 Plain Text

Requirement



pynvme builds your own tests.

Changes in SSD



- SSD has been changing for the decade:
 - media:
 - SLC, MLC, TLC, QLC
 - 2D => 3D
 - PCM, 3D-Xpoint...
 - host:
 - PATA, SATA, PCIe/NVMe
 - open-channel
 - up coming: ZNS, KV, ... ?
 - DRAM
 - form factor
- Agile: good for the constant change and uncertainty

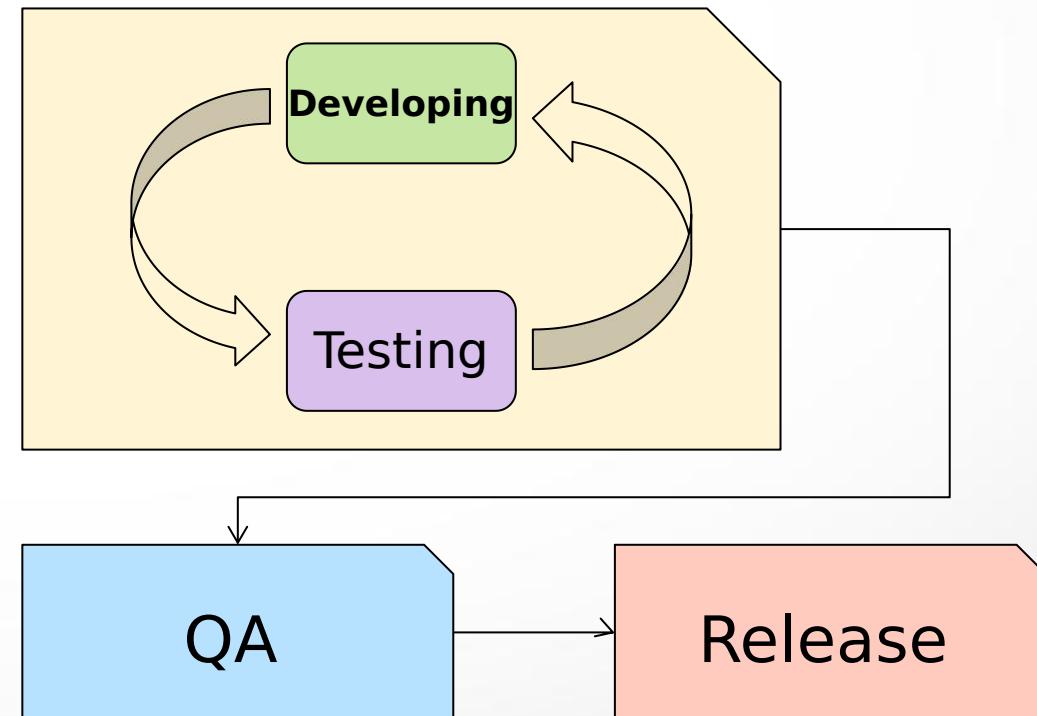


Agile Testing



- Developing and Testing are done interactively and iteratively.
- QA verifies the product of Dev/Test for customers.
- Testing and QA are different. Testing tools and QA tools are also different.
- Most available tools in the market are QA tools.

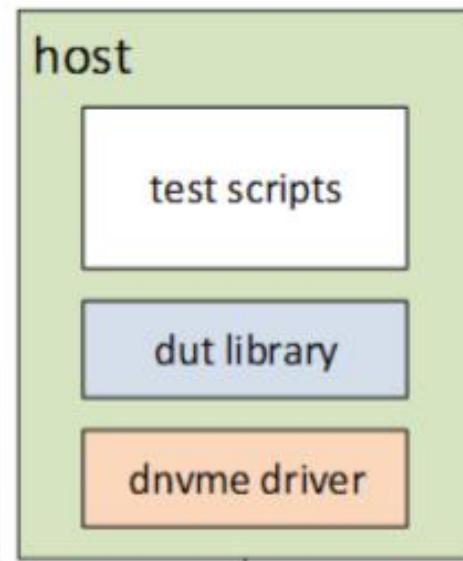
test	QA
for developer	for customer
before checkin	before release
automatic	manual
white-box	black-box
short	long
changing	stable



Experience in a Startup: dnvme



- dnvme: @2015
 - wrap with python in userspace
 - functional tests
 - integrated with Jenkins
 - firmware engineers develop scripts
 - PASS IOL test on the first try



- challenges:
 - low performance:
 - IOPS, latency, consistency
 - test efficiency
 - stress tests
 - maintainness: kernel module
 - function coverage: PRP, ...
 - GPL License
- dnvme is good at NVMe test, but in-efficient on SSD test.

Born for NVMe: SPDK



- SPDK is open from 2016
 - super high performance: [10M IOPS](#)
 - user space application
 - based on DPDK environment, which abstracts low-level resources
 - PCIe
 - memory
- meet challenges:
 - low performance:
 - IOPS, latency, consistency
 - test efficiency
 - stress tests
 - maintainness: kernel module
 - function coverage: PRP, ...
 - BSD License



SDC2020



A screenshot of a Google Chrome browser window showing the SDC 2020 website. The title bar shows multiple tabs including "SDC 2020 |". The address bar shows the URL "storagedeveloper2020.org/nvme". The page header includes the SDC 20 logo and navigation links for LOBBY, AGENDA, SPONSORS, SPEAKERS, BOFS, SMB3 IO LAB, PMEM BOOTCAMP, and CONNECT AT SDC.



pynvme: an open, fast and extensible NVMe SSD test tool

Crane Chu
GENG YUN Technology Pte Ltd

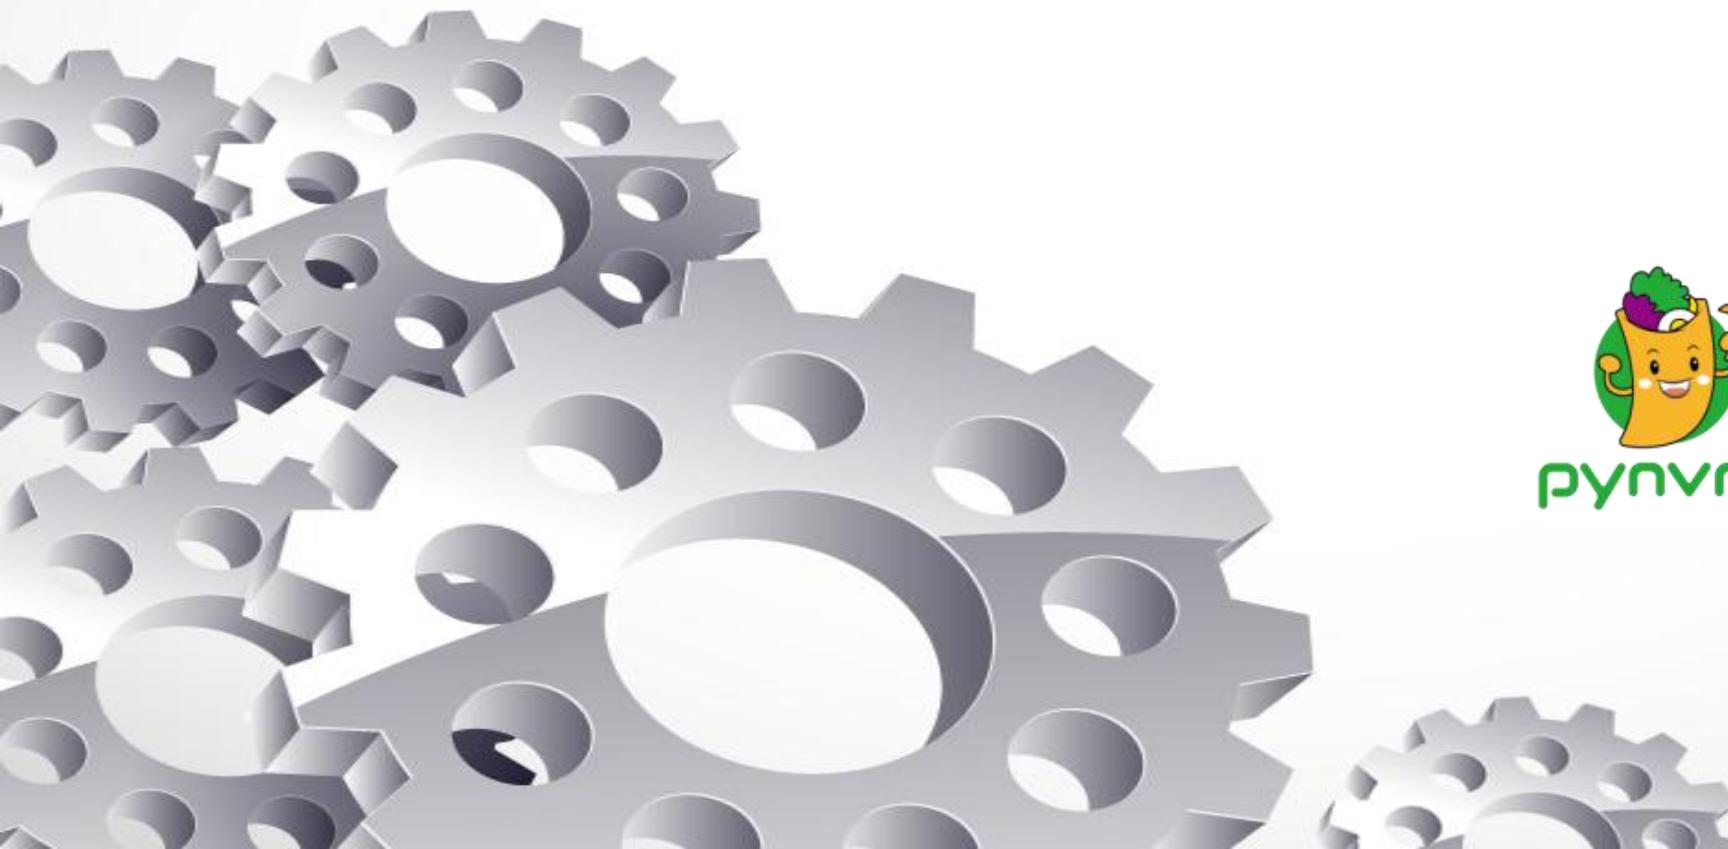
SDC 20

SDC SDC SDC
SDC SDC SDC
SDC SDC SDC
SDC SDC SDC

- Python tool for testing your NVMe design

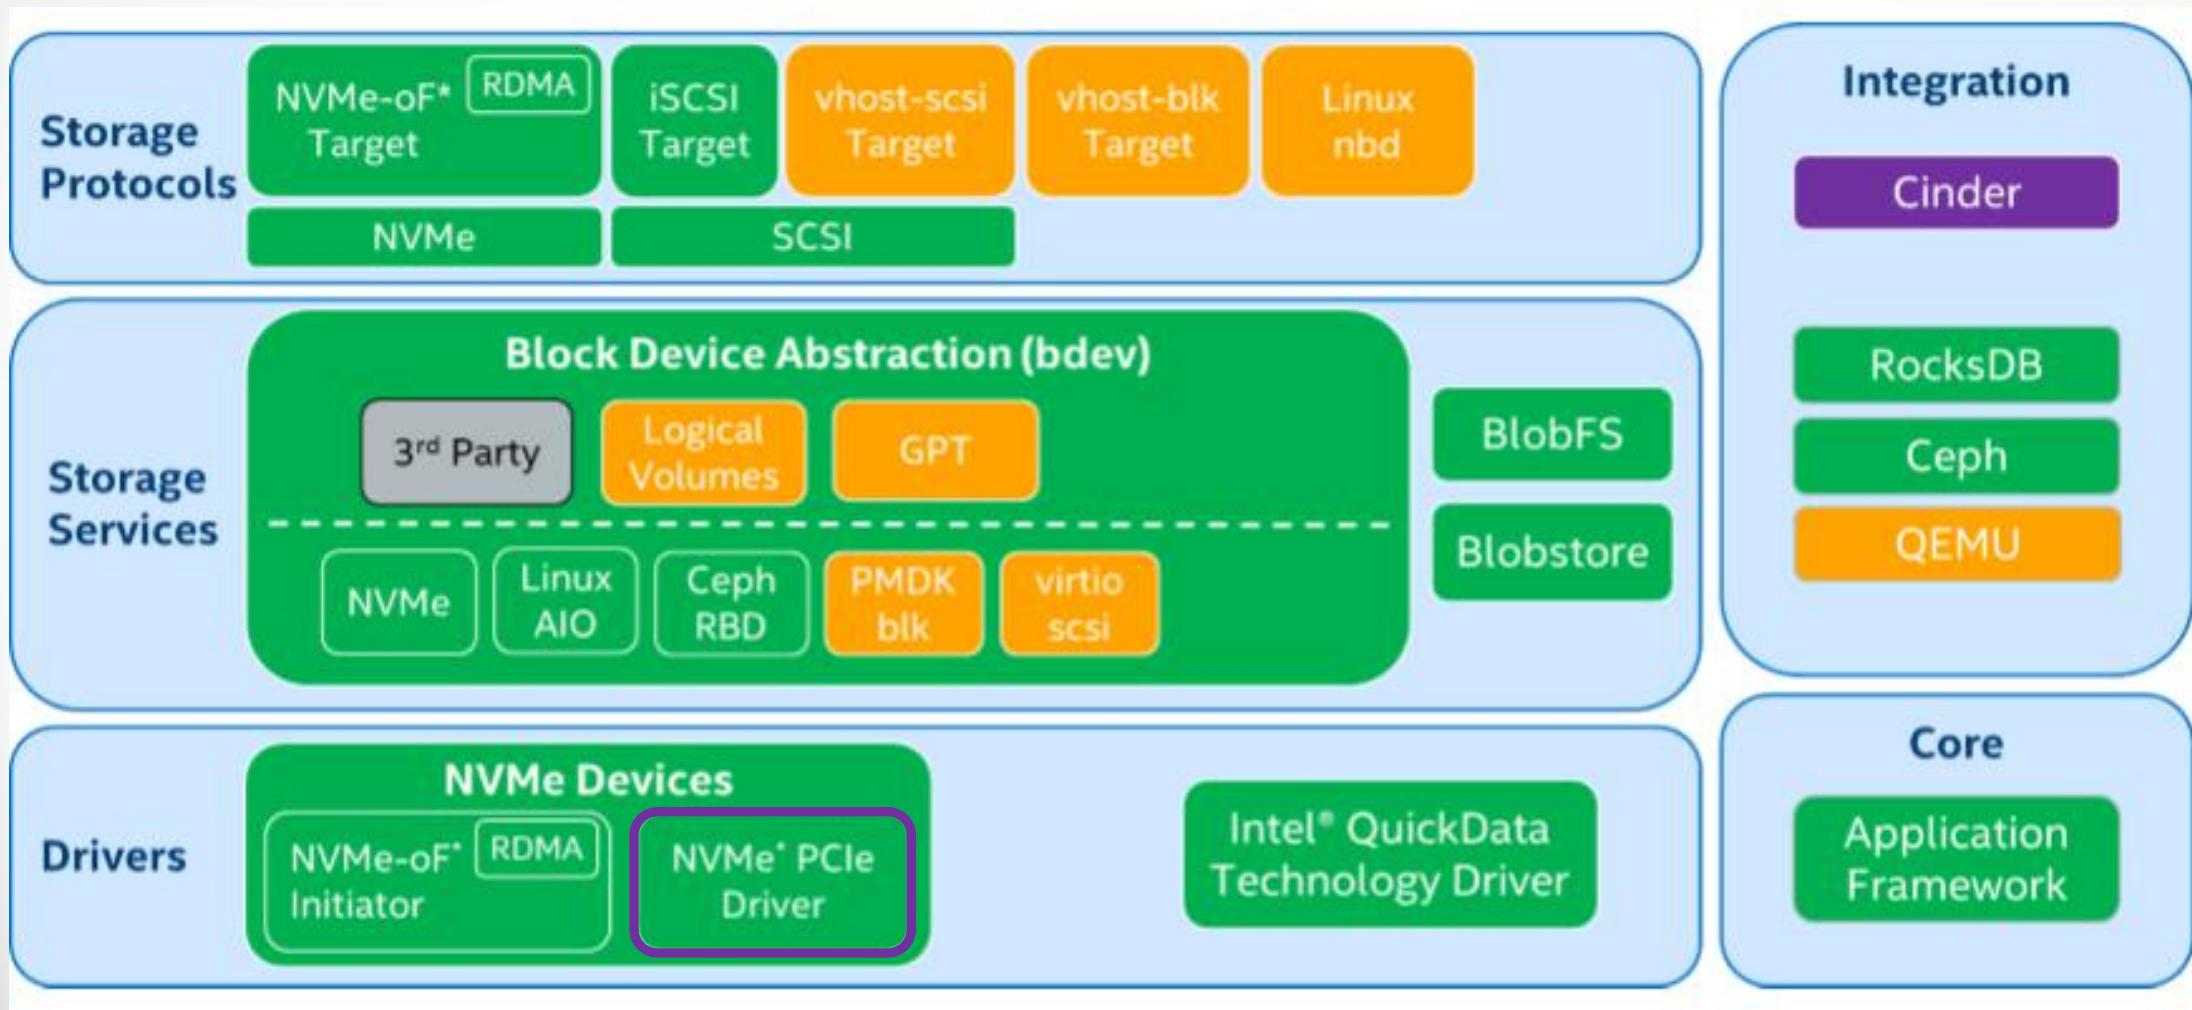
<https://www.youtube.com/watch?v=Yoru7vzVyL8>

Design



pynvme builds your own tests.

SPDK



SPDK/DPDK



- Moving all of the necessary drivers into userspace, which avoids syscalls and enables zero-copy access from the application.
- Polling hardware for completions instead of relying on interrupts, which lowers both total latency and latency variance.
- Avoiding all locks in the I/O path, instead relying on message passing.

Performance

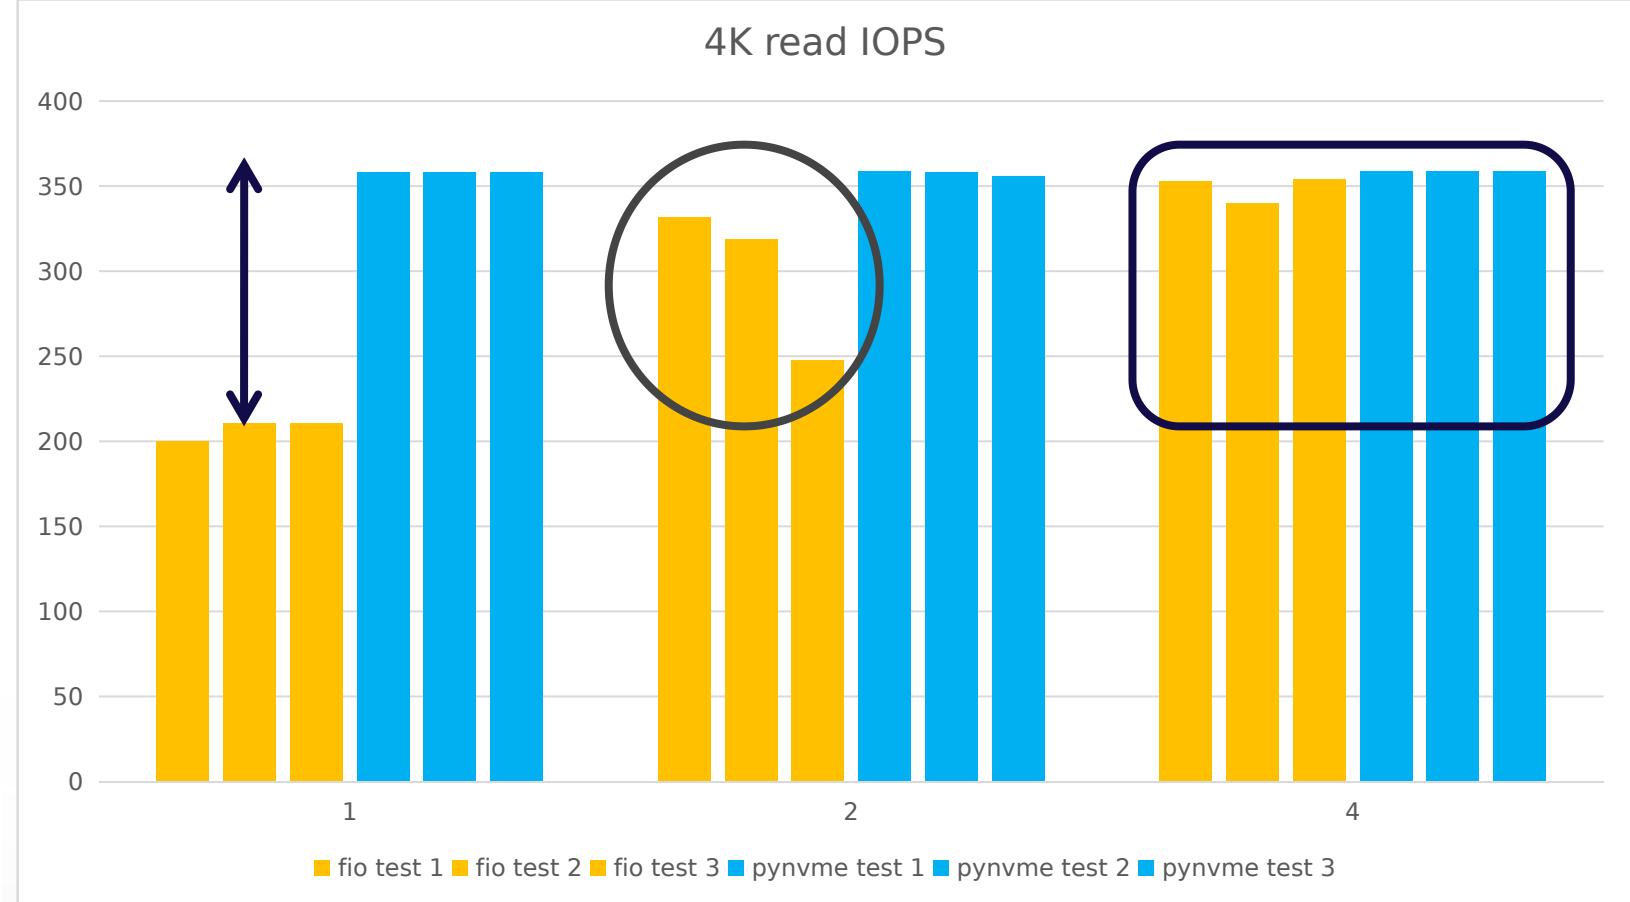


fio (unit: K IOPS)

Q count	1	2	4
test 1	200	332	353
test 2	211	319	340
test 3	211	248	354

pynvme (unit: K IOPS)

Q count	1	2	4
test 1	358	359	359
test 2	358	358	359
test 3	358	356	359



Performance: latency

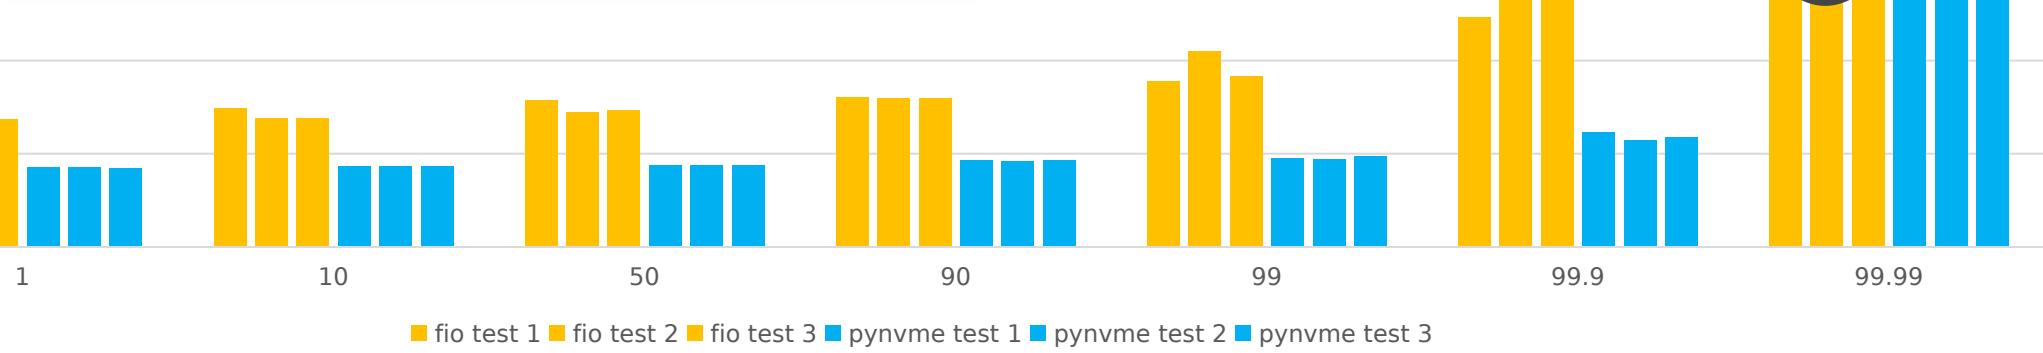


fio (unit: us)

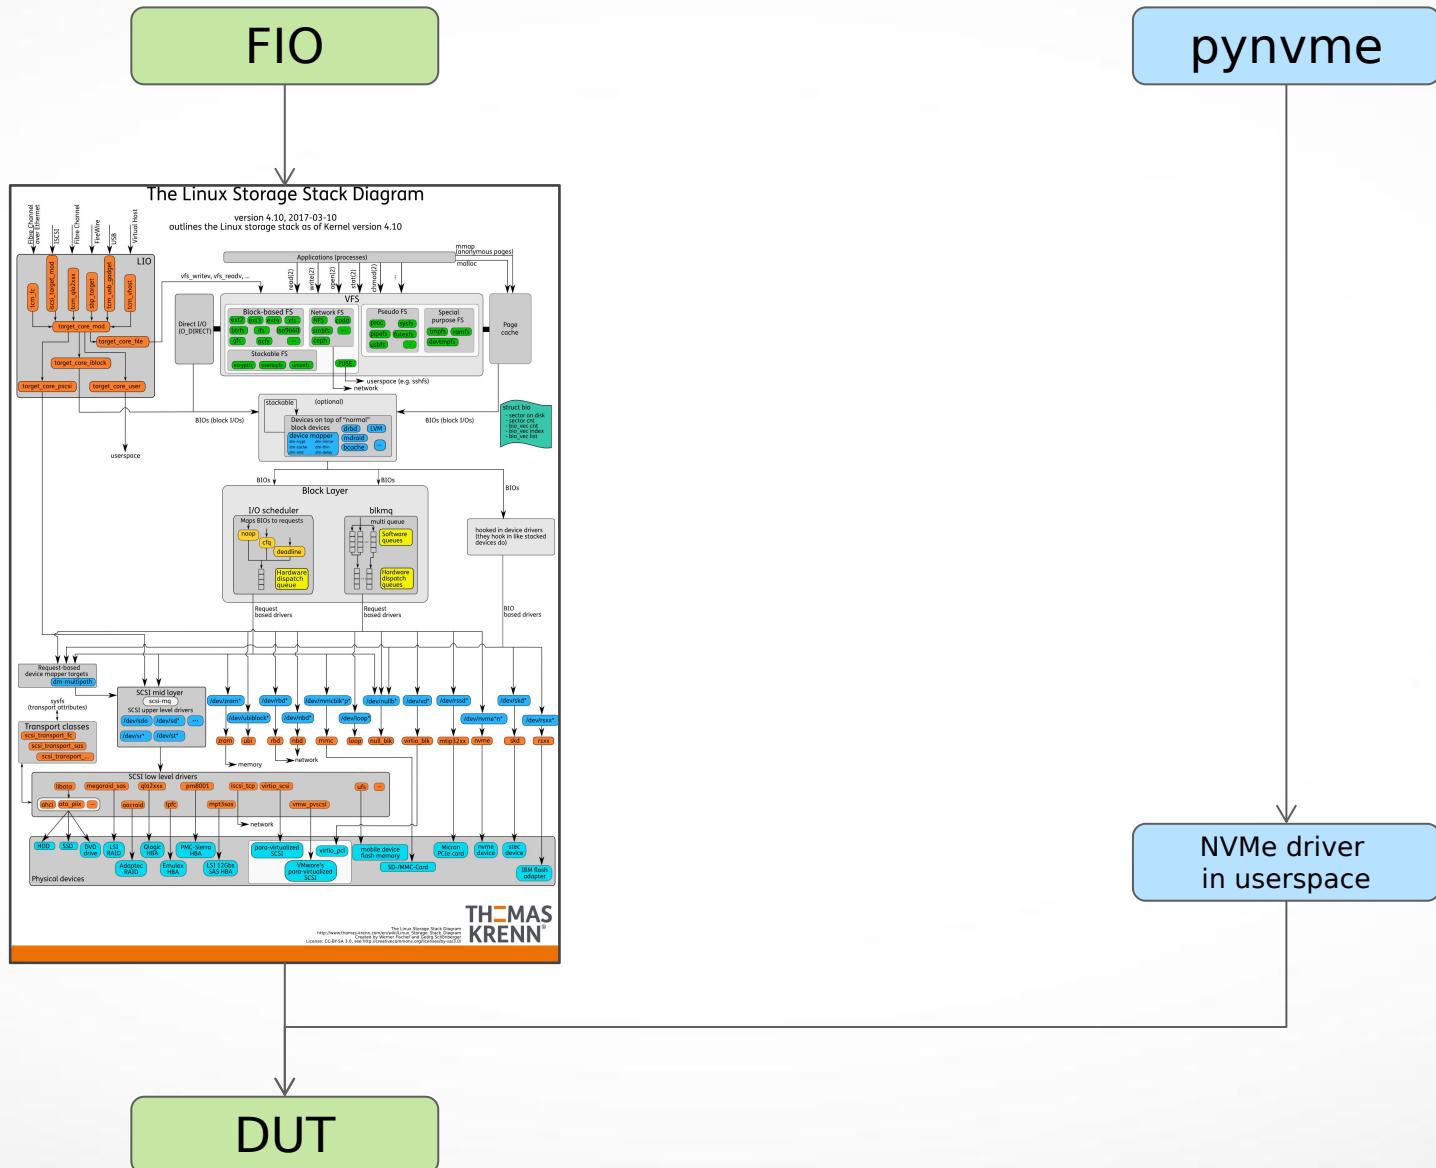
percentile	1	10	50	90	99	99.9	99.99
1	289	297	314	322	355	494	693
2	273	277	289	318	420	553	1106
3	273	277	293	318	367	644	1467

pynvme (unit: us)

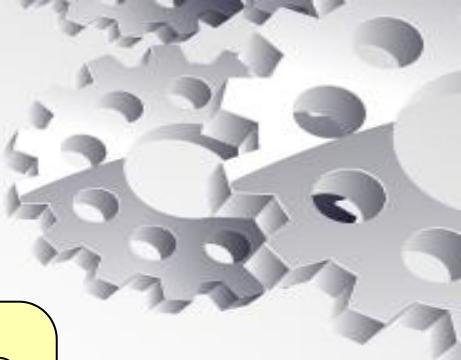
percentile	1	10	50	90	99	99.9	99.99
1	171	173	175	185	190	246	630
2	171	173	175	184	189	229	628
3	169	172	175	185	195	235	626



Performance: design



Architecture



library (ongoing):

KV

ZNS

TCG

psd

API:

controller

namespace

qpair

pcie

buffer

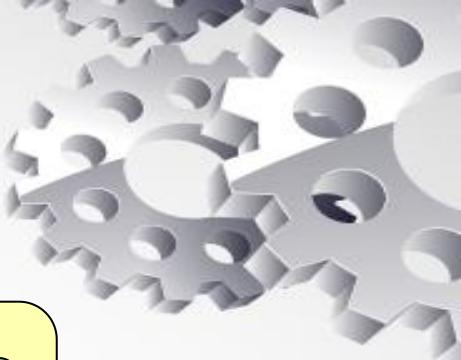
SPDK: nvme driver

cmdlog

checksum

ioworker

psd: Python Space Driver



psd: IOSQ IOCQ SQE CQE PRPList

API: controller namespace qpair pcie buffer

SPDK: nvme driver cmdlog checksum ioworker

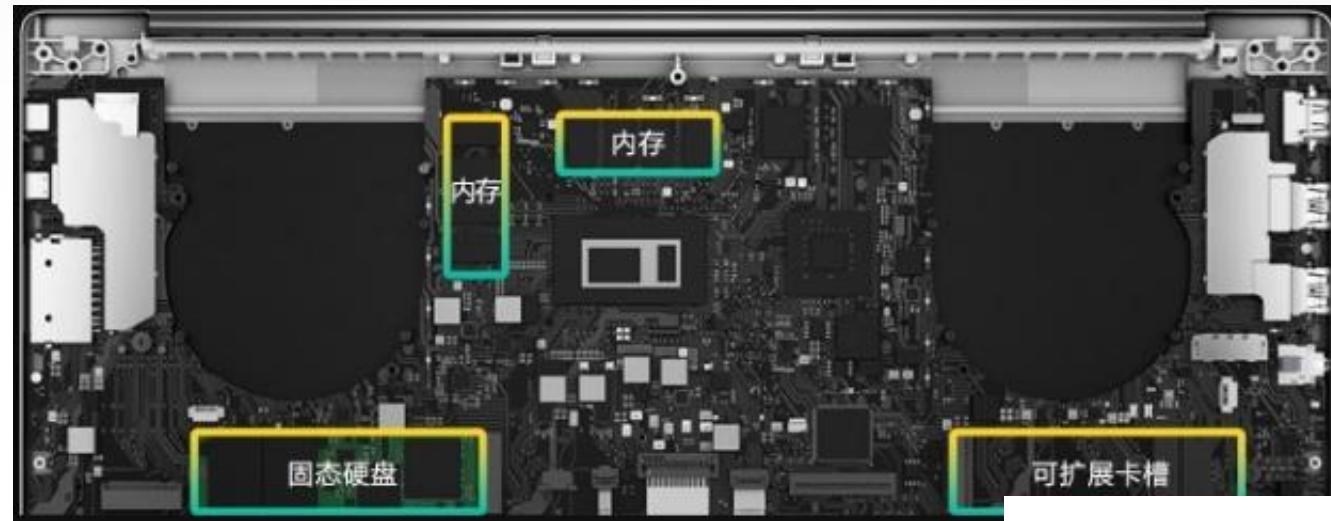
NVMe SSD

Open Ecosystem



Flexible Hardware Configuration

- Single Node
 - laptop
 - workstation
- Mass Deploy
 - server



Install and Test

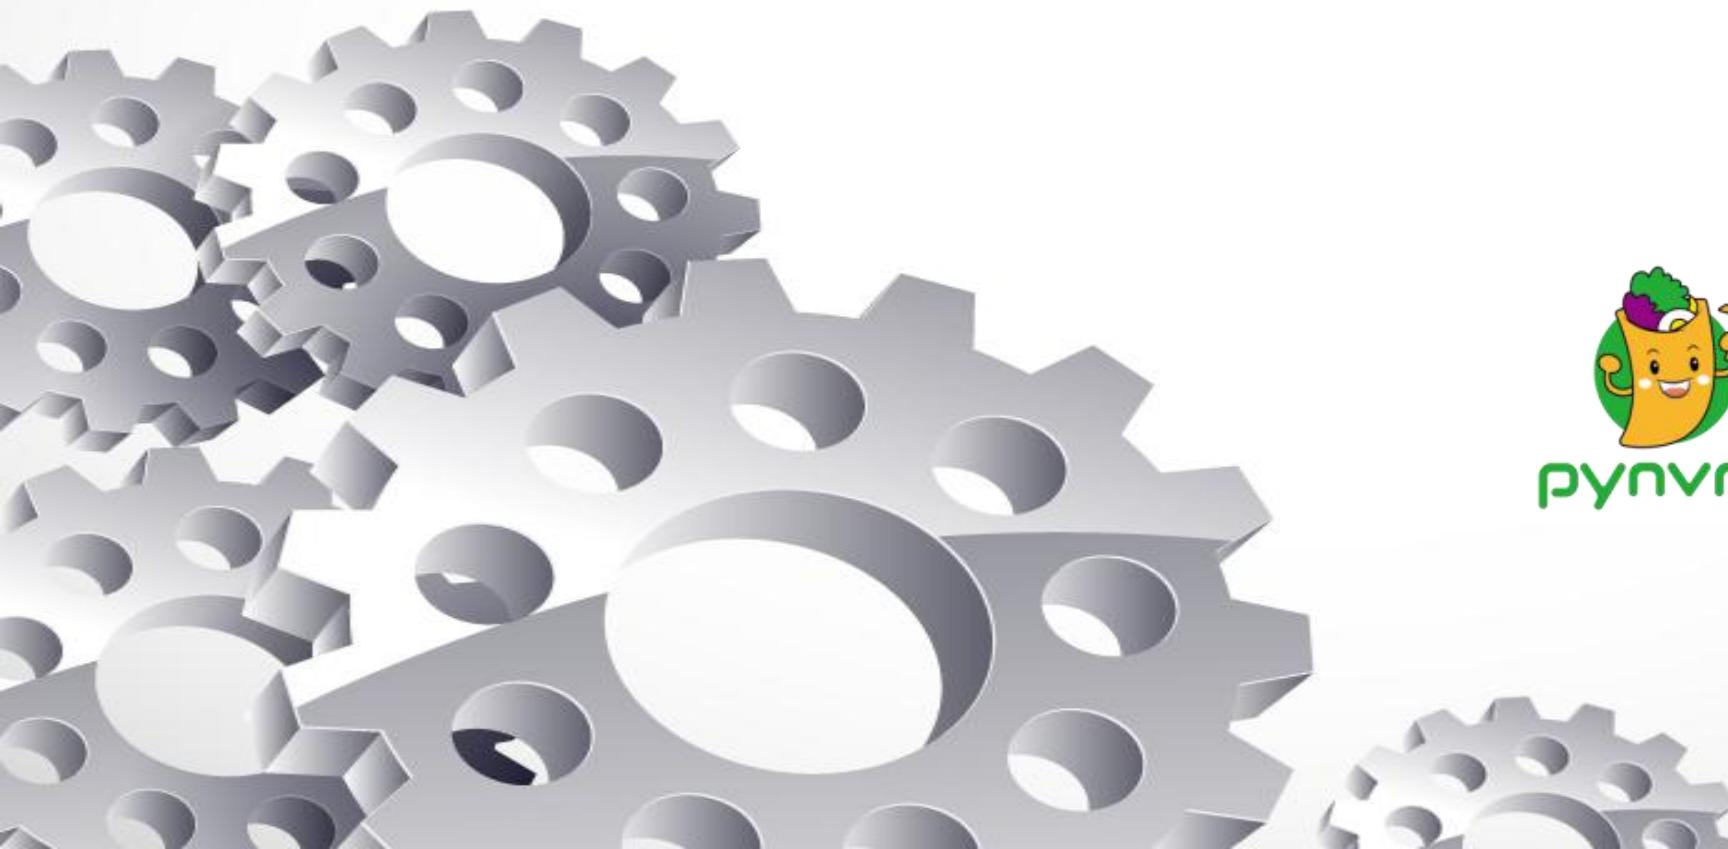


```
> git clone https://github.com/pynvme/pynvme
> cd pynvme
> ./install.sh

> make setup
> make test TESTS="driver_test.py::test_ioworker_iops_multiple_queue[1]"

> make test
TESTS=scripts/stress/endurance_test.py::test_replay_jedec_client_trace
pciaddr=02:00.0
> make test
TESTS=scripts/stress/endurance_test.py::test_replay_jedec_client_trace
pciaddr=172.168.5.44
```

Scripts



pynvme builds your own tests.

Test Scripts

```
1 import time
2 import pytest
3 import logging
4
5 import nvme as d
6
7
8 # intuitive, spec, qpair, vscode, debug, cmdlog, assert
9 def test_hello_world(nvme0, nvme0n1, qpair):
10     # prepare data buffer and IO queue
11     read_buf = d.Buffer(512)
12     write_buf = d.Buffer(512)
13     write_buf[10:21] = b'hello world'
14
15     # send write and read command
16     def write_cb(cdw0, status1): # command callback function
17         nvme0n1.read(qpair, read_buf, 0, 1)
18         nvme0n1.write(qpair, write_buf, 0, 1, cb=write_cb)
19
20     # wait commands complete and verify data
21     assert read_buf[10:21] != b'hello world'
22     qpair.waitdone(2)
23     assert read_buf[10:21] == b'hello world'
```

```
8     def test_quarch_dirty_power_cycle_single(nvme0, nvme0n1, subsystem, buf, verify):
9         # get the unsafe shutdown count before test
10        nvme0.getlogpage(2, buf, 512).waitdone()
11        orig_unsafe_count = buf.data(159, 144)
12        logging.info("unsafe shutdowns: %d" % orig_unsafe_count)
13        assert verify == True
14
15        # 128K random write
16        cmdlog_list = [None]*1000
17        with nvme0n1.ioworker(io_size=256,
18                               lba_random=True,
19                               read_percentage=30,
20                               region_end=256*1000*1000,
21                               time=30,
22                               qdepth=1024,
23                               output_cmdlog_list=cmdlog_list):
24            # sudden power loss before the ioworker end
25            time.sleep(10)
26            subsystem.poweroff()
27
28            # power on and reset controller
29            time.sleep(5)
30            subsystem.poweron()
31            time.sleep(0)
32            nvme0.reset()
```

Test Scripts: dirty power cycle



- [PCIe Card Module](#)
- [Torridon Interface Kit](#)
- poweroff process
 - poweroff when ioworker is alive
 - remove device from system
- poweron process
 - poweron
 - remove kernel driver
 - rescan device
 - nvme initialization



Features



- access PCI configuration space
- access NVMe registers in BAR space
- send any NVMe admin/IO commands
- support callback functions for NVMe commands
- support MSI/MSIx interrupts
- transparent checksum verification on every LBA
- generates IO workload of high performance and low latency
- support multiple namespaces
- support multiple tests on different controllers
- integrate with the test framework pytest
- integrate with VSCode to display cmdlog in GUI
- support NVMe over TCP targets
- doc: <https://pynvme.readthedocs.io/>

Test Scripts: 3 ways of sending IO



```
8 # intuitive, spec, qpair, vscode, debug, cmdlog, assert
9 def test_hello_world(nvme0, nvme0n1, qpair):
10     # prepare data buffer and IO queue
11     read_buf = d.Buffer(512)
12     write_buf = d.Buffer(512)
13     write_buf[10:21] = b'hello world'
14
15     # send write and read command
16     def write_cb(cdw0, status1): # command callback function
17         nvme0n1.read(qpair, read_buf, 0, 1)
18         nvme0n1.write(qpair, write_buf, 0, 1, cb=write_cb)
```

```
def test_ioworker_simplified(nvme0n1):
    nvme0n1.ioworker(io_size=2, time=2).start().close()
```

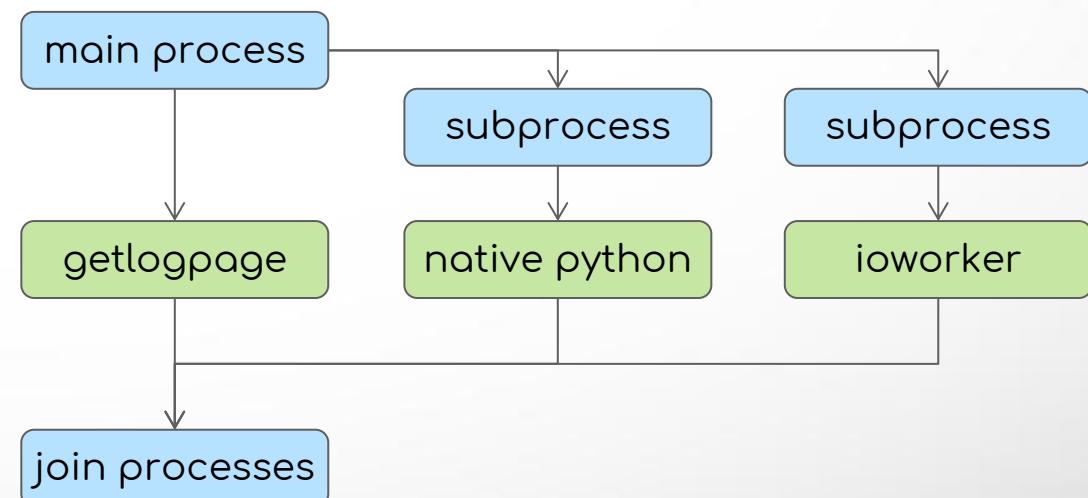
```
698 def test_write_before_power_cycle(nvme0, subsystem):
699     cq = IOCQ(nvme0, 1, 128, PRP(2*1024))
700     sq = IOSQ(nvme0, 1, 128, PRP(8*1024), cqid=1)
701
702     #burst write
703     for i in range(127):
704         cmd = SQE(1, 1)
705         buf = PRP(512, ptype=32, pvalue=i)
706         cmd.prp1 = buf
707         cmd[10] = i
708         sq[i] = cmd
709
710     # write 127 512byte at one shot
711     sq.tail = 127
```

Test Scripts: multiprocessing



```
72 def test_ioworker_with_temperature_and_trim(nvme0, nvme0n1):
73     # start trim process
74     import multiprocessing
75     mp = multiprocessing.get_context("spawn")
76     p = mp.Process(target = subprocess Trim,
77                     args = (nvme0.addr.encode('utf-8'),
78                             300000))
79     p.start()
80
81     # start read/write ioworker and admin commands
82     smart_log = d.Buffer(512, "smart log")
83     with nvme0n1.ioworker(io_size=8, lba_align=16,
84                           lba_random=True, qdepth=16,
85                           read_percentage=67, iops=10000, time=10):
86         for i in range(15):
87             nvme0.getlogpage(0x02, smart_log, 512).waitdone()
88             ktemp = smart_log.data(2, 1)
89
90             from pytemperature import k2c
91             logging.info("temperature: %0.2f degreeC" % k2c(ktemp))
92             time.sleep(1)
93
94     # wait trim process complete
95     p.join()
96
```

```
60     # ioworker with admin commands, multiprocessing, log, cmdlog, pythonic
61     def subprocess Trim(pciaddr, loops):
62         nvme0 = d.Controller(pciaddr)
63         nvme0n1 = d.Namespace(nvme0)
64         q = d.Qpair(nvme0, 8)
65         buf = d.Buffer(4096)
66         buf.set_dsm_range(0, 8, 8)
67
68         # send trim commands
69         for i in range(loops):
70             nvme0n1.dsm(q, buf, 1).waitdone()
71
```



Test Scripts: power related

```
188 # PCIe, different of power states and resets
189 def test_power_and_reset(pcie, nvme0, subsystem):
190     pcie.aspm = 2                      # ASPM L1
191     pcie.power_state = 3                # PCI PM D3hot
192     pcie.aspm = 0
193     pcie.power_state = 0
194
195     nvme0.reset()                     # controller reset: CC.EN
196     nvme0.getfeatures(7).waitdone()
197
198     pcie.reset()                      # PCIe reset: hot reset, TS1, TS2
199     nvme0.reset()                     # reset controller after pcie reset
200     nvme0.getfeatures(7).waitdone()
201
202     pcie.flr()                        # PCIe function level reset
203     nvme0.reset()                     # reset controller after pcie reset
204     nvme0.getfeatures(7).waitdone()
205
206     subsystem.reset()                 # NVMe subsystem reset: NSSR
207     nvme0.reset()                     # controller reset: CC.EN
208     nvme0.getfeatures(7).waitdone()
209
210     subsystem.power_cycle(10)        # power cycle NVMe device: cold reset
211     nvme0.reset()                     # controller reset: CC.EN
212     nvme0.getfeatures(7).waitdone()
```

```
42 @pytest.mark.parametrize("ps", [4, 3, 2, 1, 0])
43 def test_format_at_power_state(nvme0, nvme0n1, ps):
44     nvme0.setfeatures(0x2, cdw11=ps).waitdone()
45     assert nvme0n1.format(ses=0) == 0
46     assert nvme0n1.format(ses=1) == 0
47     p = nvme0.getfeatures(0x2).waitdone()
48     assert p == ps
```

```
191 @pytest.mark.parametrize("aspm", [0, 2])
192 @pytest.mark.parametrize("gen", [1, 2, 3, 2, 1, 1, 2, 3])
193 def test_PCIE_link_speed(pcie, nvme0, nvme0n1, gen, aspm):
194     linkctr2_addr = pcie.cap_offset(0x10)+0x30
195     linkctr2 = pcie.register(linkctr2_addr, 4)
196     logging.info(linkctr2)
197
198     pcie[linkctr2_addr] = (linkctr2 & 0xf0) | gen
199     logging.info(pcie.register(linkctr2_addr, 4))
200     pcie.reset()
201     nvme0.reset()
202     test_PCIE_read_bandwidth(nvme0n1)
203     test_PCIE_link_control_aspm(nvme0, pcie, aspm)
```

Test Scripts: NVMe initialization for WRR

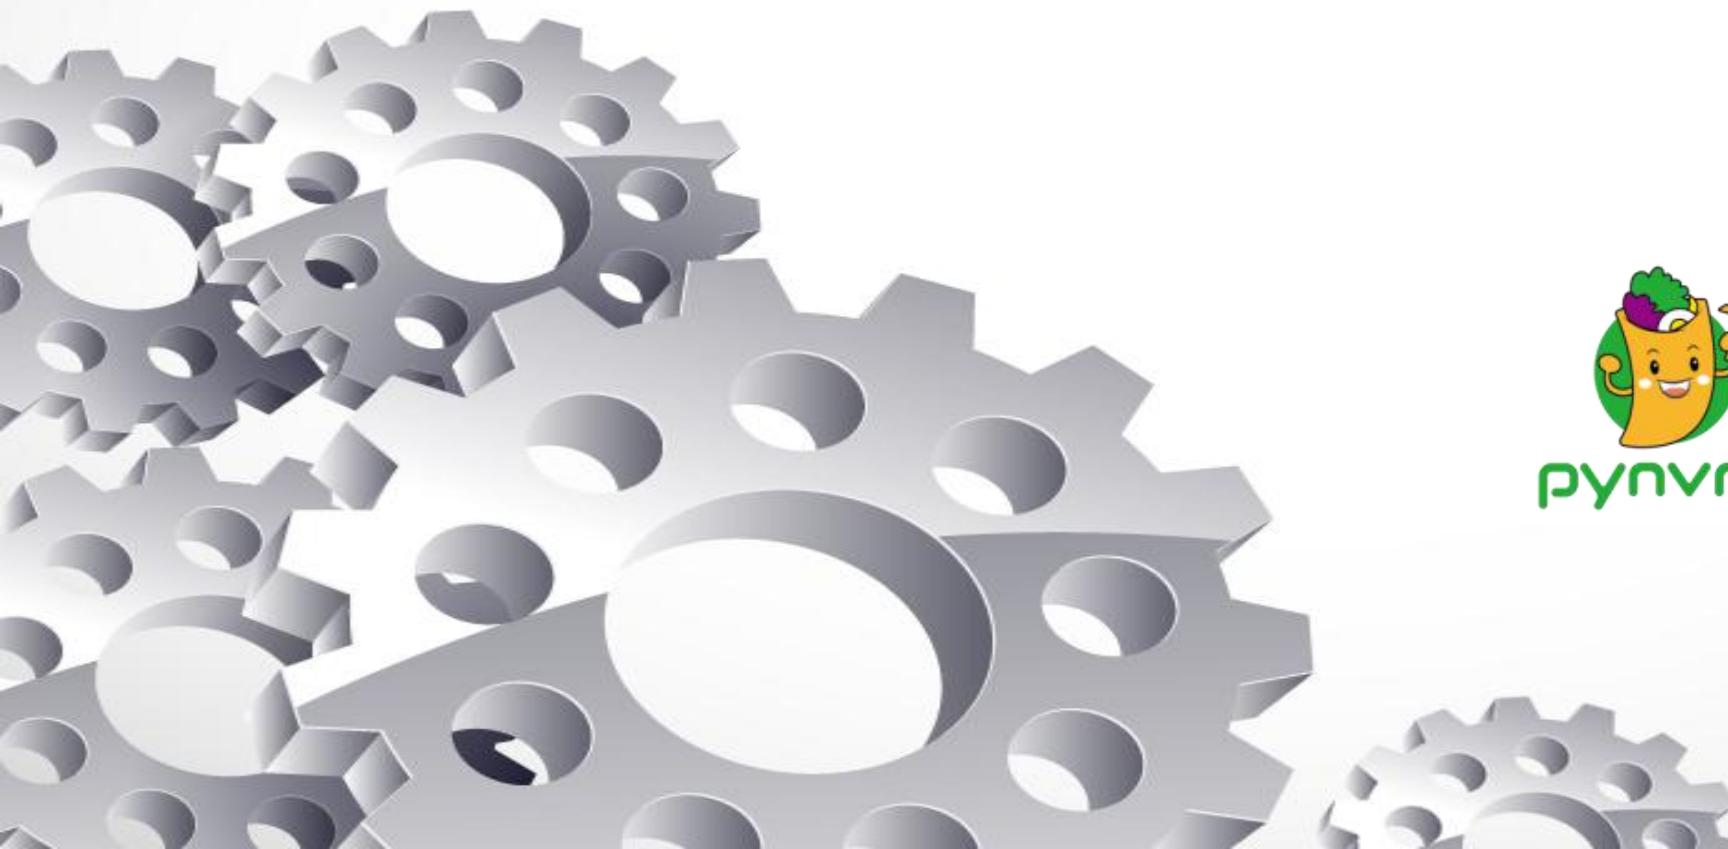
```
9  def nvme_init_wrr(nvme0):
10     logging.info("user defined nvme init")
11
12     nvme0[0x14] = 0
13     while not (nvme0[0x1c]&0x1) == 0: pass
14
15     # 3. set admin queue registers
16     nvme0.init_adminq()
17
18     # 4. set register cc
19     if (nvme0.cap>>17) & 0x1:
20         logging.info("set WRR arbitration")
21         nvme0[0x14] = 0x00460800
22     else:
23         nvme0[0x14] = 0x00460000
24
25     # 5. enable cc.en
26     nvme0[0x14] = nvme0[0x14] | 1
27
28     # 6. wait csts.rdy to 1
29     while not (nvme0[0x1c]&0x1) == 1: pass
30
31     # 7. identify controller
32     nvme0.identify(Buffer(4096)).waitdone()
33
34     # 8. create and identify all namespace
35     nvme0.init_ns()
36
37     # 9. set/get num of queues
38     nvme0.setfeatures(0x7, cdw11=0x00ff00ff).waitdone()
39     nvme0.getfeatures(0x7).waitdone()
```

Test Scripts: error code



```
93 def test_write_in_sanitize_operations(nvme0, nvme0n1, buf, qpairs):
94     if nvme0.id_data(331, 328) == 0: #L9
95         pytest.skip("sanitize operation is not supported") #L10
96
97     logging.info("supported sanitize operation: %d" % nvme0.id_data(331, 328))
98     nvme0.sanitize().waitdone() #L13
99
100    with pytest.warns(UserWarning, match="ERROR status: 00/1d"):
101        nvme0n1.write(qpair, buf, 0).waitdone()
102
103    # check sanitize status in log page
104    with pytest.warns(UserWarning, match="AER notification is triggered"):
105        nvme0.getlogpage(0x81, buf, 20).waitdone() #L17
106        while buf.data(3, 2) & 0x7 != 1: #L18
107            time.sleep(1)
108            nvme0.getlogpage(0x81, buf, 20).waitdone() #L20
109            progress = buf.data(1, 0)*100//0xffff
110            logging.info("%d%%" % progress)
111
112    # check sanitize status
113    nvme0.getlogpage(0x81, buf, 20).waitdone()
114    assert buf.data(3, 2) & 0x7 == 1
115
```

ioworker



pynvme builds your own tests.

ioworker: Parameters



Input Parameters:

- lba_start, lba_step, lba_align
- lba_random
- region_start, region_end
- distribution
- io_size (int, range, list, dict)
- read_percentage, **op_percentage**
- time, io_count
- qdepth, qprio
- pvalue, ptype
- **iops**
- ...

Output Parameters:

- io_count_read
- io_count_write, io_count_nonread
- mseconds
- latency_max_us
- latency_average_us
- error
- cpu_usage
- output_io_per_second (optional)
- output_percentile_latency (optional)
- output_cmdlog_list (optional)
- ...

API document: <https://pynvme.readthedocs.io/api.html>

ioworker: Python API

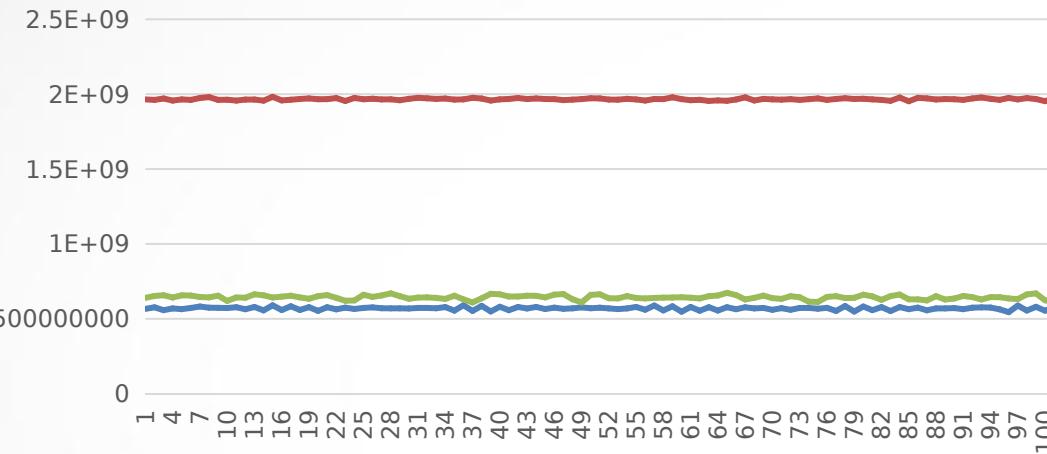


```
774 def test_ioworker_with_temperature(nvme0, nvme0n1, buf):
775     with nvme0n1.ioworker(io_size=256,
776                           time=30,
777                           op_percentage={0:10, # flush
778                                         2:60, # read
779                                         9:30}), \
780         nvme0n1.ioworker(io_size=8,
781                           time=30,
782                           op_percentage={0:10, # flush
783                                         9:10, # trim
784                                         1:80}):# write
785     for i in range(40):
786         time.sleep(1)
787         nvme0.getlogpage(0x02, buf, 512).waitdone()
788         ktemp = buf.data(2, 1)
789         from pytemperature import k2c
790         logging.info("temperature: %0.2f degreec" %
791                      k2c(ktemp))
792
```

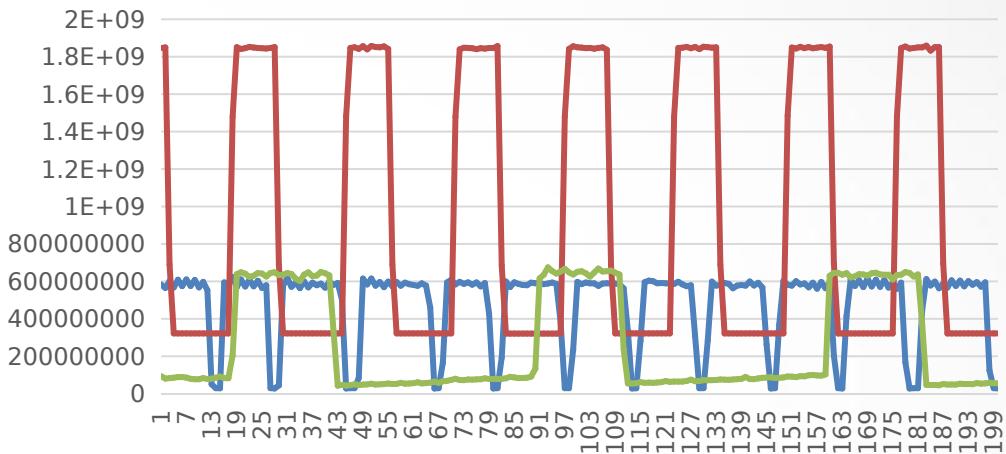
ioworker: Gallery



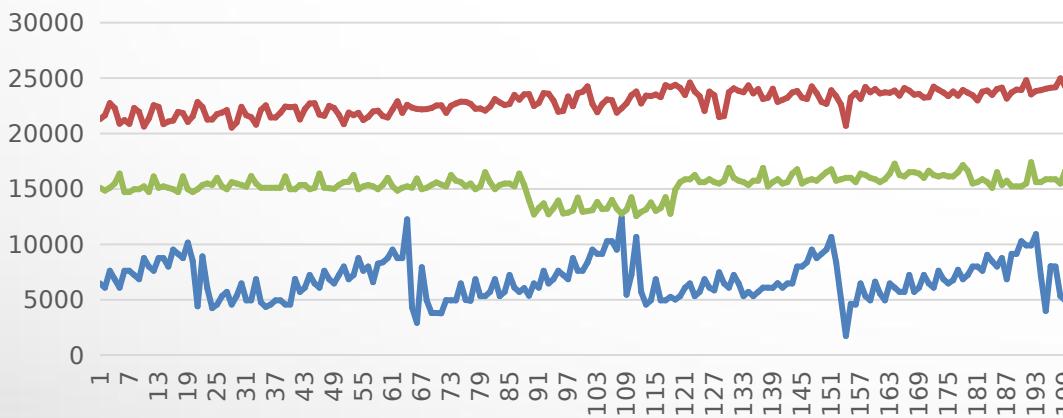
sequential write 1st pass (x: second, y: B/s)



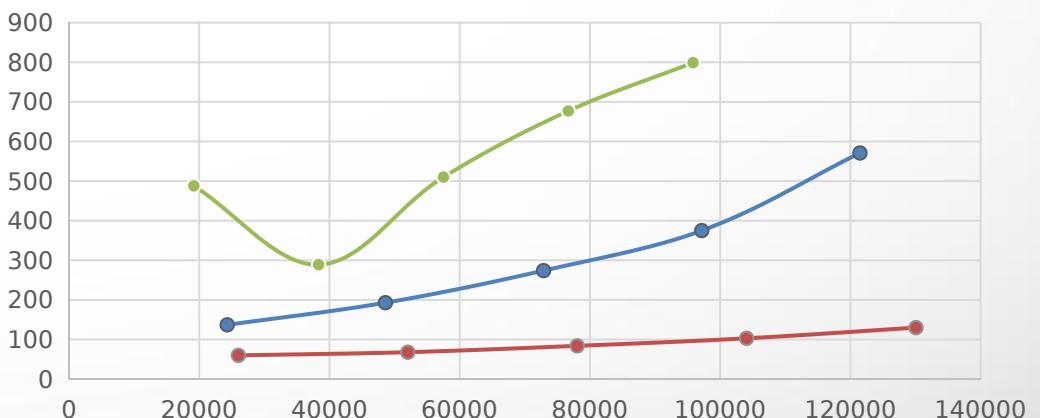
sequential write 3rd pass (x: second, y: B/s)



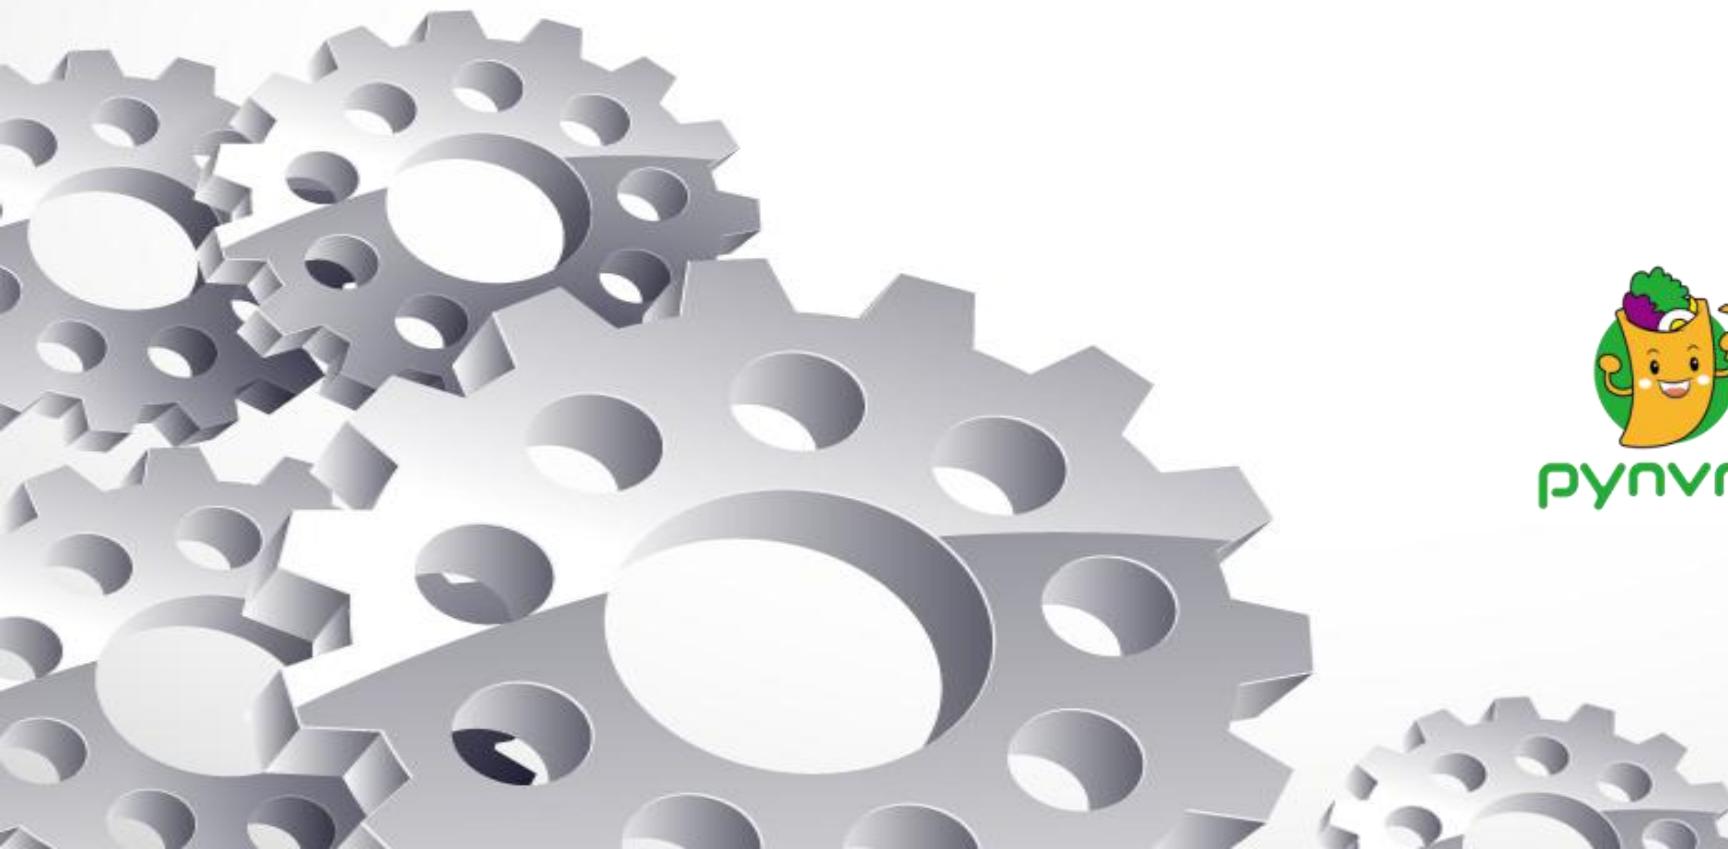
4K random write (x: second, y: IOPS)



latency against IOPS, 2R1W (x: IOPS, y: us)



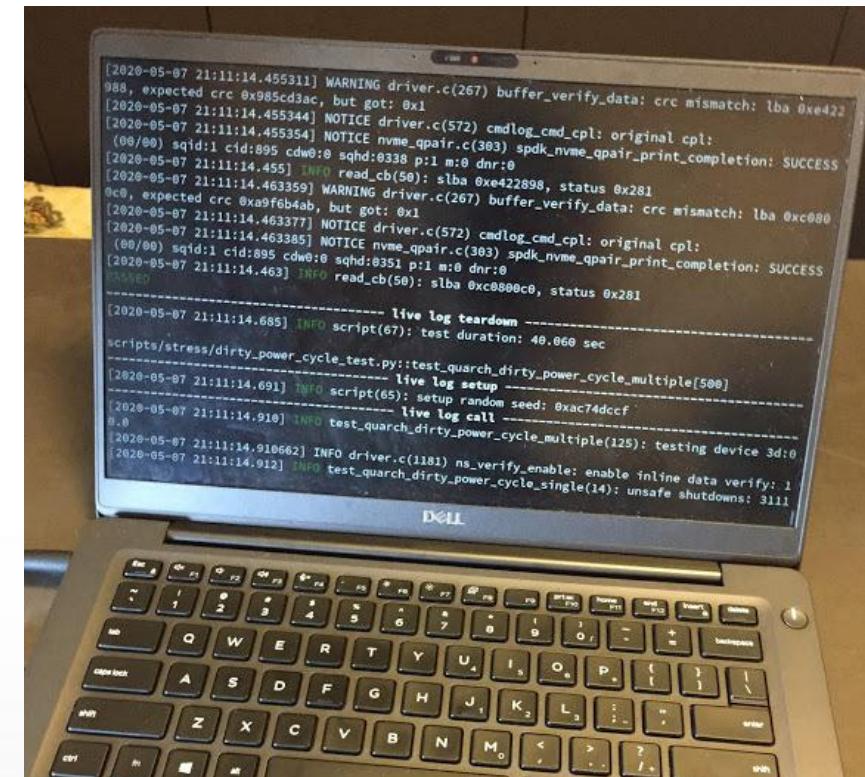
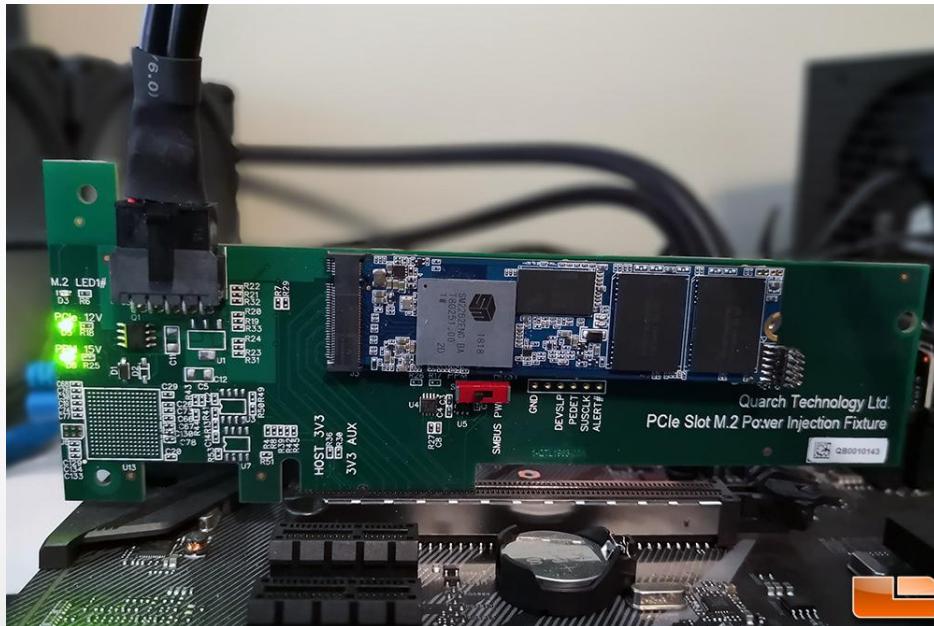
Features



pynvme builds your own tests.

vPower to ON and OFF

- Usually we need a special hardware box to control the power of SSD
- vPower controls the power of SSD:
 - power off: enter S3/sleep mode
 - power on: wake up by RTC



vAnalyzer

TEST

PYNVME QPAIRS

0000:3d:00.0 Q00: >>

0000:3d:00.0 Q01: >>

0000:3d:00.0 Q02: >>

0000:3d:00.0 Q03: >>

0000:3d:00.0 Q04: >>

0000:3d:00.0 Q05: >>

0000:3d:00.0 Q06: >>

0000:3d:00.0 Q07: >>

0000:3d:00.0 Q08: >>

0000:3d:00.0 Q09: >> **selected**

0000:3d:00.0 Q10: >>

0000:3d:00.0 Q11: >>

0000:3d:00.0 Q12: >>

0000:3d:00.0 Q13: >>

0000:3d:00.0 Q14: >>

0000:3d:00.0 Q15: >>

0000:3d:00.0 Q16: >>

PYTHON

scripts

conformance

ftl

performance

snia

stress

tcg

trace

test_examples.py

test_utilities.py

LOG 0000:3d:00.0 Q02

CMDLOG 0000:3d:00.0 Q05

G 0000:3d:00.0 Q02

CMDLOG 0000:3d:00.0 Q13

CMDLOG 0000:3d:00.0 Q00

Performance Gauge

1353.80 MB/s

330.52 K/s

1 2020-04-15 13:33:21.463617 [cmd001: Read]

2 0x00050002, 0x00000001, 0x00000000, 0x00000000

3 0x00000000, 0x00000000, 0x0c7d3000, 0x00000000

4 0x00000000, 0x00000000, 0x213d9b90, 0x00000000

5 0x00000007, 0x00000000, 0x00000000, 0x00000000

6 not completed

7 ...

9 2020-04-15 13:33:21.463613 [cmd002: Read]

10 0x00010002, 0x00000001, 0x00000000, 0x00000000

11 0x00000000, 0x00000000, 0x0c7d7000, 0x00000000

12 0x00000000, 0x00000000, 0x16b26aa0, 0x00000000

13 0x00000007, 0x00000000, 0x00000000, 0x00000000

14 not completed

15 ...

17 2020-04-15 13:33:21.463607 [cmd003: Read]

18 0x00030002, 0x00000001, 0x00000000, 0x00000000

19 0x00000000, 0x00000000, 0x0c7db000, 0x00000000

20 0x00000000, 0x00000000, 0x151c50c0, 0x00000000

21 0x00000007, 0x00000000, 0x00000000, 0x00000000

22 not completed

1 2020-04-15 13:32:39.104405 [cmd001: Create I/O Submission Queue]

2 0x005a0001, 0x00000000, 0x00000000, 0x00000000

3 0x00000000, 0x00000000, 0x0ea0000, 0x00000001

4 0x00000000, 0x00000000, 0x00070010, 0x00100001

5 0x00000000, 0x00000000, 0x00000000, 0x00000000

6 2020-04-15 13:32:39.104519: [cpl: SUCCESS]

7 0x00000000, 0x00000000, 0x00000056, 0x0000005a

8 ...

9 2020-04-15 13:32:39.103458 [cmd002: Create I/O Completion Queue]

10 0x005a0005, 0x00000000, 0x00000000, 0x00000000

11 0x00000000, 0x00000000, 0x0e800000, 0x00000001

12 0x00000000, 0x00000000, 0x00070010, 0x00100003

13 0x00000000, 0x00000000, 0x00000000, 0x00000000

14 2020-04-15 13:32:39.104403: [cpl: SUCCESS]

15 0x00000000, 0x00000000, 0x00000055, 0x0000005a

16 ...

17 2020-04-15 13:32:38.533985 [cmd003: Create I/O Submission Queue]

18 0x005a0001, 0x00000000, 0x00000000, 0x00000000

19 0x00000000, 0x00000000, 0x0e600000, 0x00000001

20 0x00000000, 0x00000000, 0x0007000f, 0x000f0001

21 0x00000000, 0x00000000, 0x00000000, 0x00000000

22 2020-04-15 13:32:38.534648: [cpl: SUCCESS]

23 0x00000000, 0x00000000, 0x00000054, 0x0000005a

24 ...

25 2020-04-15 13:32:38.533838 [cmd004: Create I/O Completion Queue]

26 0x005a0005, 0x00000000, 0x00000000, 0x00000000

27 0x00000000, 0x00000000, 0x0e400000, 0x00000001

28 0x00000000, 0x00000000, 0x0007000f, 0x000f0003

29 0x00000000, 0x00000000, 0x00000000, 0x00000000

30 2020-04-15 13:32:38.533979: [cpl: SUCCESS]

31 0x00000000, 0x00000000, 0x00000053, 0x0000005a

32 ...

33 2020-04-15 13:32:38.313247 [cmd005: Create I/O Submission Queue]

34 0x005a0001, 0x00000000, 0x00000000, 0x00000000

35 0x00000000, 0x00000000, 0x0e200000, 0x00000001

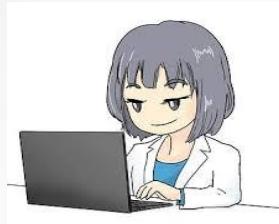
36 0x00000000, 0x00000000, 0x0007000e, 0x000e0001

37 0x00000000, 0x00000000, 0x00000000, 0x00000000

master* Python 3.7.6 64-bit 0 △ 0 Run Tests

Ln 1, Col 1 Spaces: 4 Plain Text

vTracer: IO Recorder and Replayer



Windows 10 Virtual Machine

QEMU with pynvme vSSD

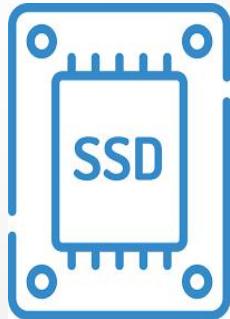
Linux

record



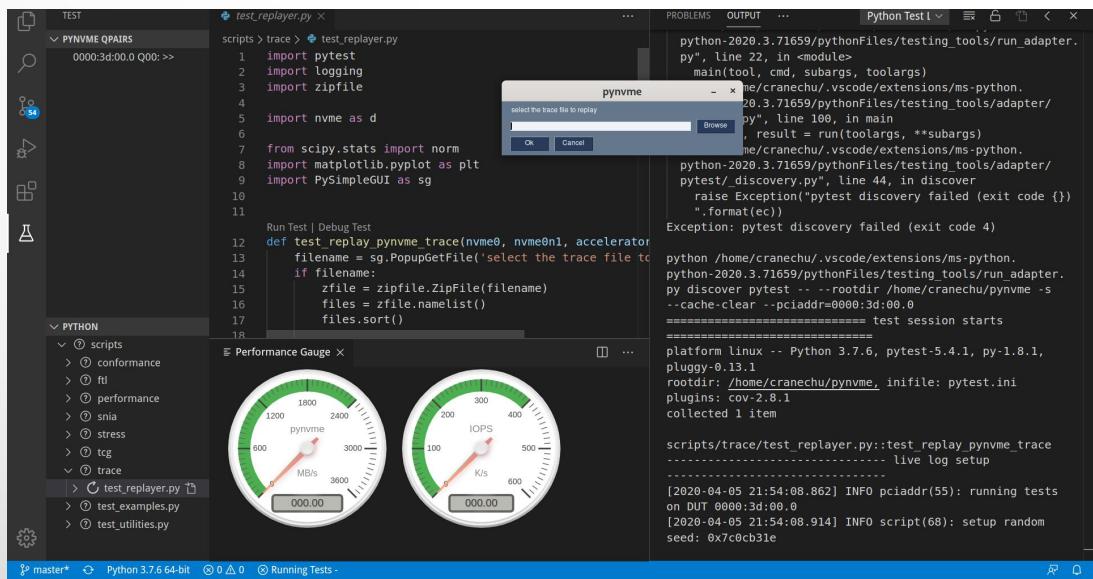
replay

pynvme ioworker



vTracer: Demo

```
Laptop:/pynvme/scripts/trace> ./home/cranechu/gemu/build/x86_64-softmmu/qemu-system-x86_64 -m 4096 -drive file=/home/cranechu/vm/win10_100g.qcow2,lf=none,id=drv0,discard=on -device nvme,drive=drv0,foo -device usb-ehci,id=ehci -device usb-tablet,bus=ehci.0 -enable-kvm -cpu host -smp 4,sockets=1,cores=4,threads=1 | ./recorder.py  
trace time base 0  
create minute trace file /tmp/pynvme_trace/0/1  
create minute trace file /tmp/pynvme_trace/0/2  
create minute trace file /tmp/pynvme_trace/0/3  
create minute trace file /tmp/pynvme_trace/0/4  
  
  
qemu-system-x86_64: warning: guest updated active QH  
  
qemu-system-x86_64: warning: guest updated active QH
```

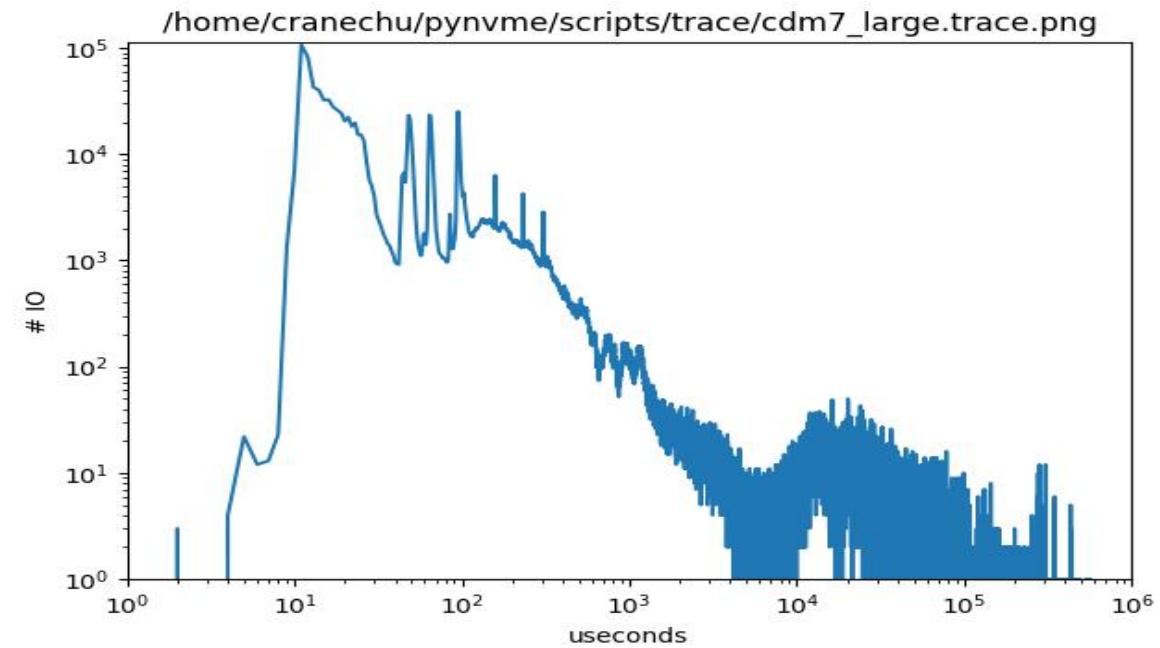
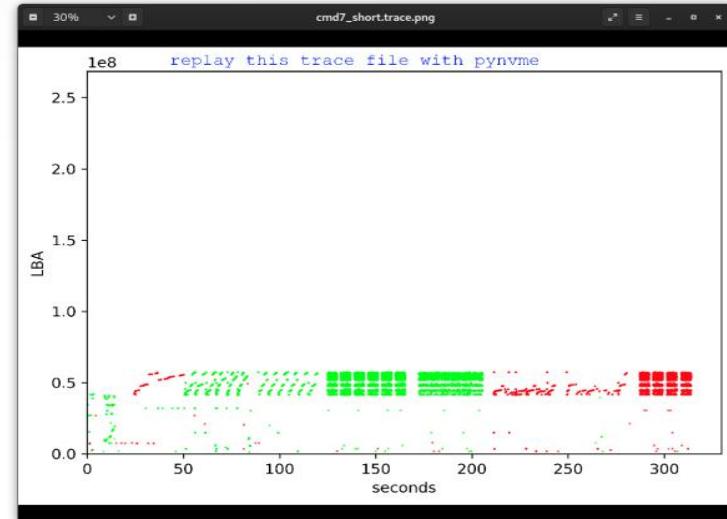


TEST PYNVME QPAIRS 0000:3d:00.0 Q0:00: > test_replayer.py
scripts > trace > test_replayer.py
0000:3d:00.0 Q0:00: >
1 import pytest
2 import logging
3 import zipfile
4
5 import nvme as d
6
7 from scipy.stats import norm
8 import matplotlib.pyplot as plt
9 import PysimpleGUI as sg
10
11 Run Test | Debug Test
12 def test_replay_pynvme_trace(nvme0, nvme0n1, accelerator
13 filename = sg.PopupGetFile('select the trace file to
14 if filename:
15 zfile = zipfile.ZipFile(filename)
16 files = zfile.namelist()
17 files.sort()
18
19
PROBLEMS OUTPUT ... Python Testl < x
pynvme select the trace file to replay
...
python-2020.3.71659/pythonFiles/testing_tools/run_adapter.py", line 22, in <module>
main(tool, cmd, subargs, toolargs)
- x me/cranechu/.vscode/extensions/ms-python.
20.3.71659/pythonFiles/testing_tools/adapter/
py", line 100, in main
, result = run(toolargs, **subargs)
me/cranechu/.vscode/extensions/ms-python.
python-2020.3.71659/pythonFiles/testing_tools/adapter/
pytest/discovery.py", line 44, in discover
raise Exception("pytest discovery failed (exit code {})".format(ec))
Exception: pytest discovery failed (exit code 4)
python /home/cranechu/.vscode/extensions/ms-python.
python-2020.3.71659/pythonFiles/testing_tools/run_adapter.
py discover pytest --rootdir /home/cranechu/pynvme -s
--cache-clear --pciaid=0000:3d:00.0
===== test session starts
platform linux -- Python 3.7.6, pytest-5.4.1, py-1.8.1,
pluggy-0.13.1
rootdir: /home/cranechu/pynvme, inifile: pytest.ini
plugins: cov-2.8.1
collected 1 item

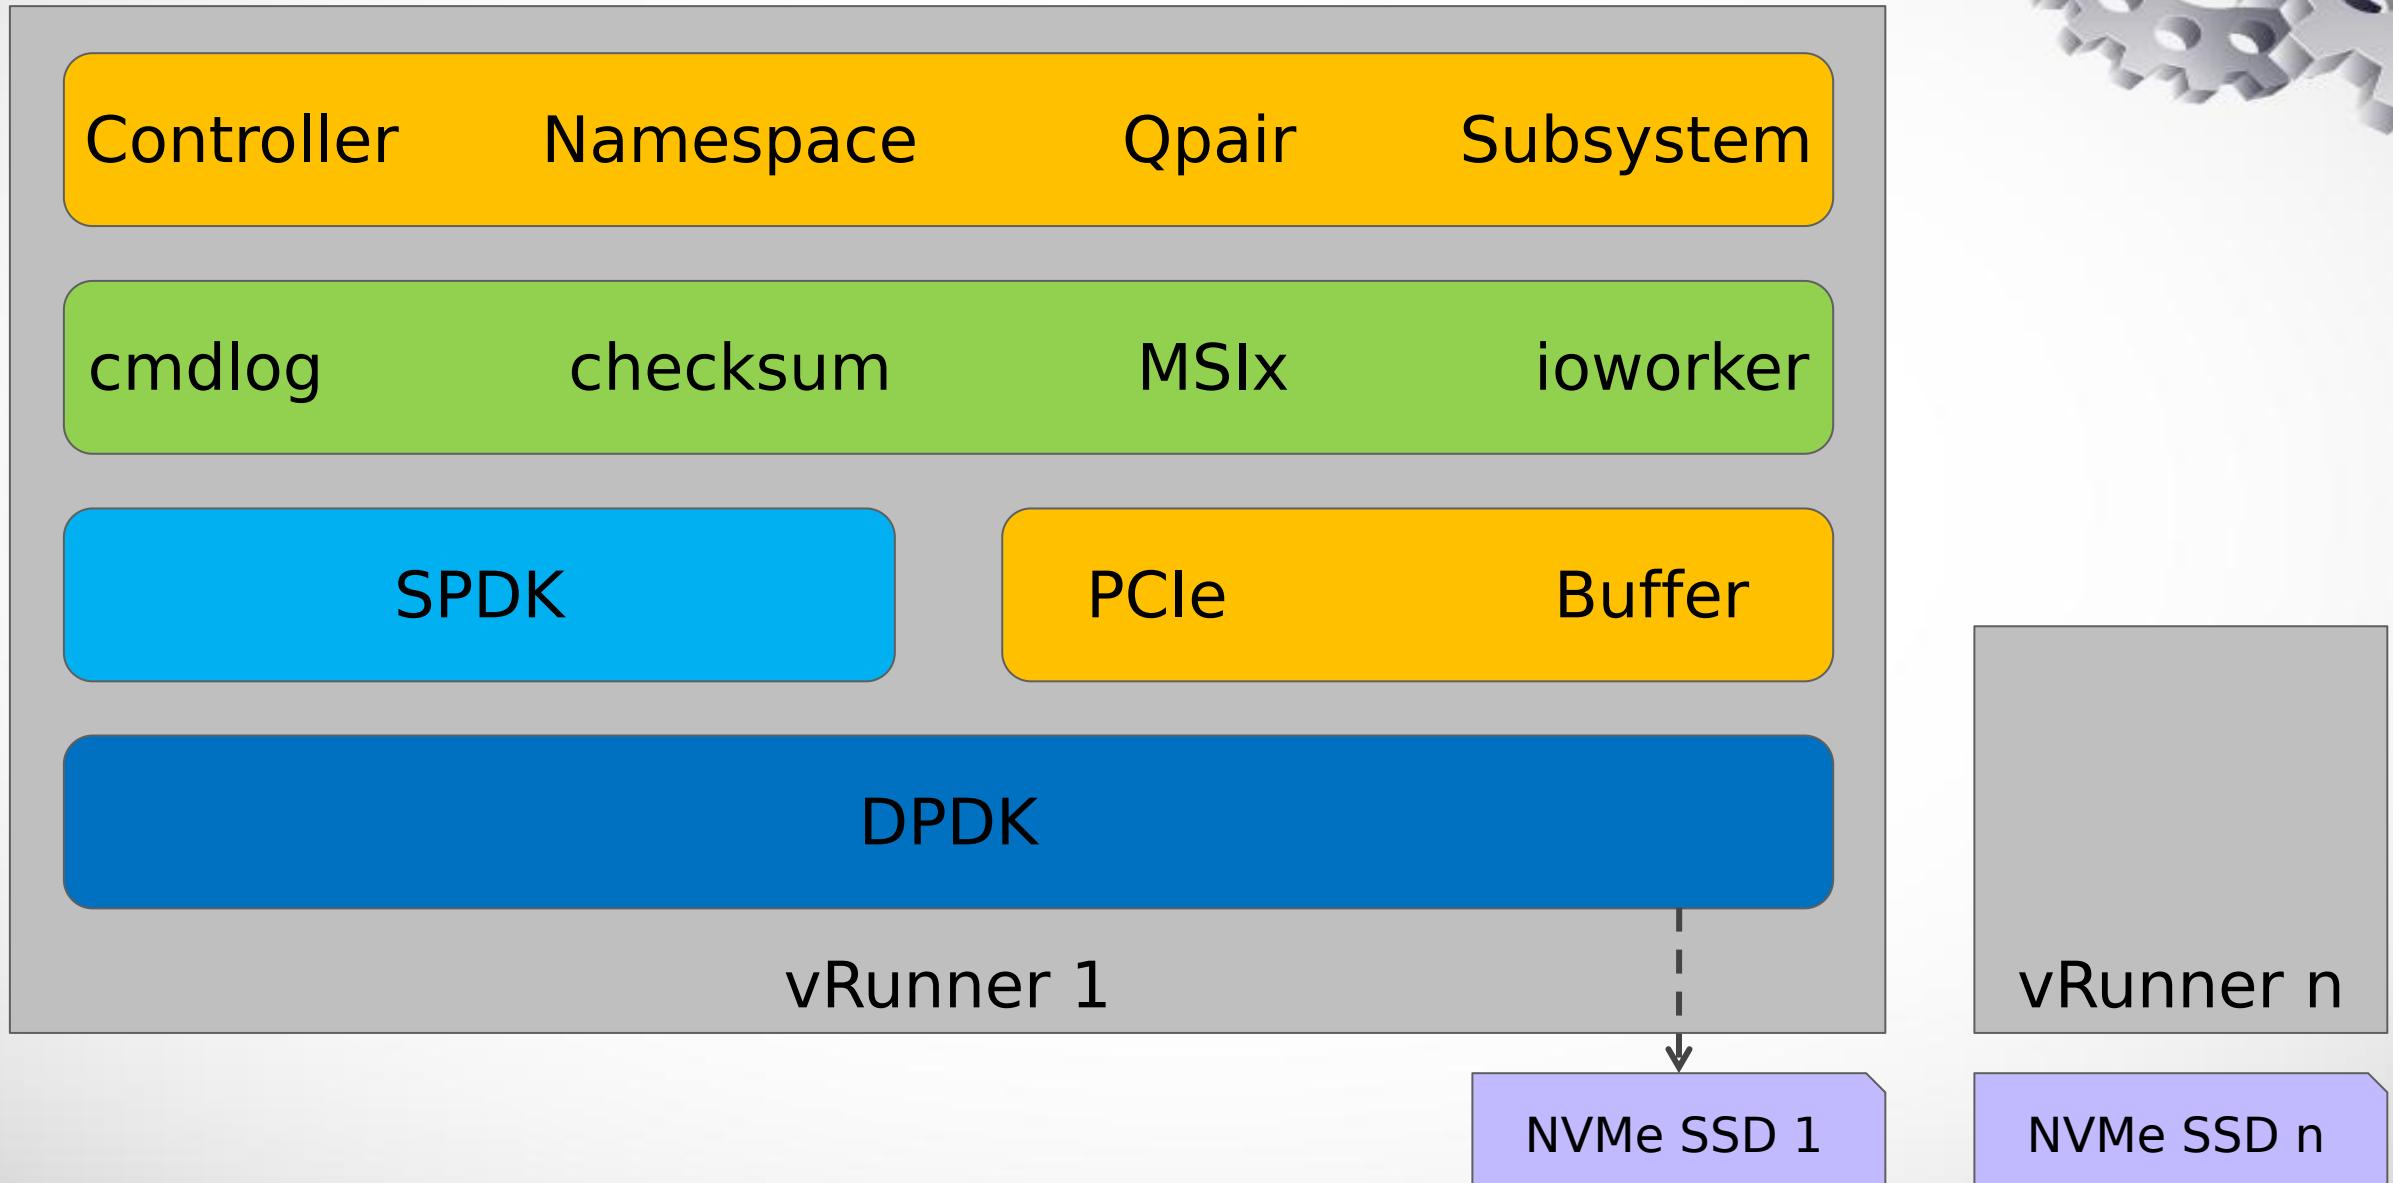
scripts/trace/test_replayer.py::test_replay_pynvme_trace
----- live log setup
[2020-04-05 21:54:08.862] INFO pciaddr(55): running tests
on DUT 0000:3d:00.0
[2020-04-05 21:54:08.914] INFO script(68): setup random
seed: 0x7c0cb31e

Performance Gauge X
pynvme MB/s 000.00 3600
IOPS K/s 000.00 600

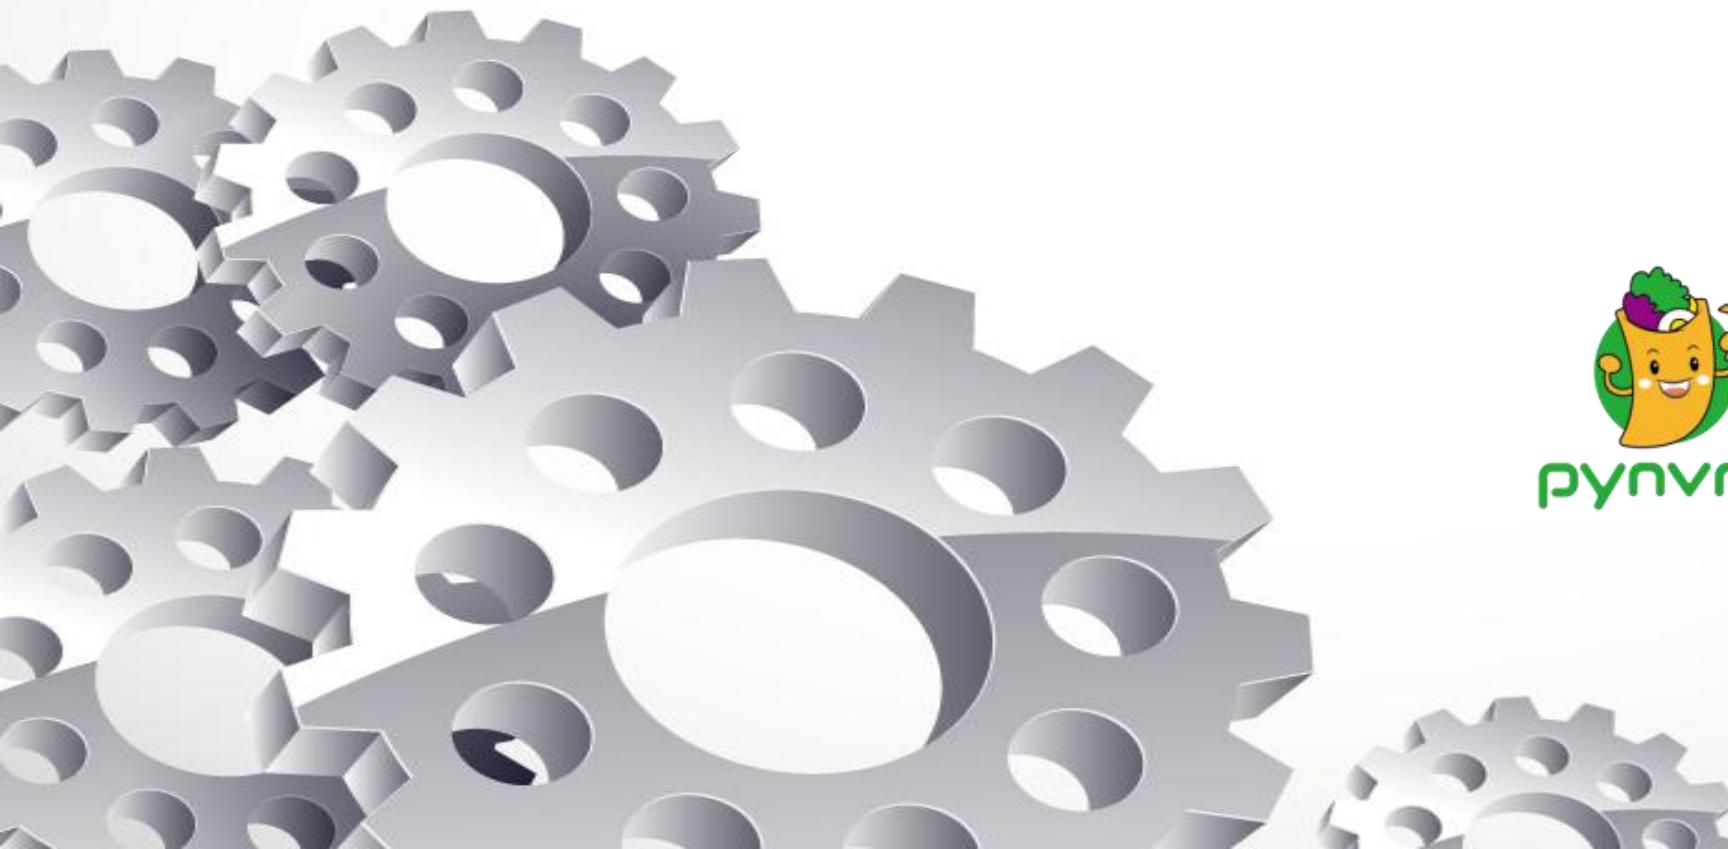
master* Python 3.7.6 64-bit 0 0 0 Running Tests -



vRunner: Mass Deploy



Services



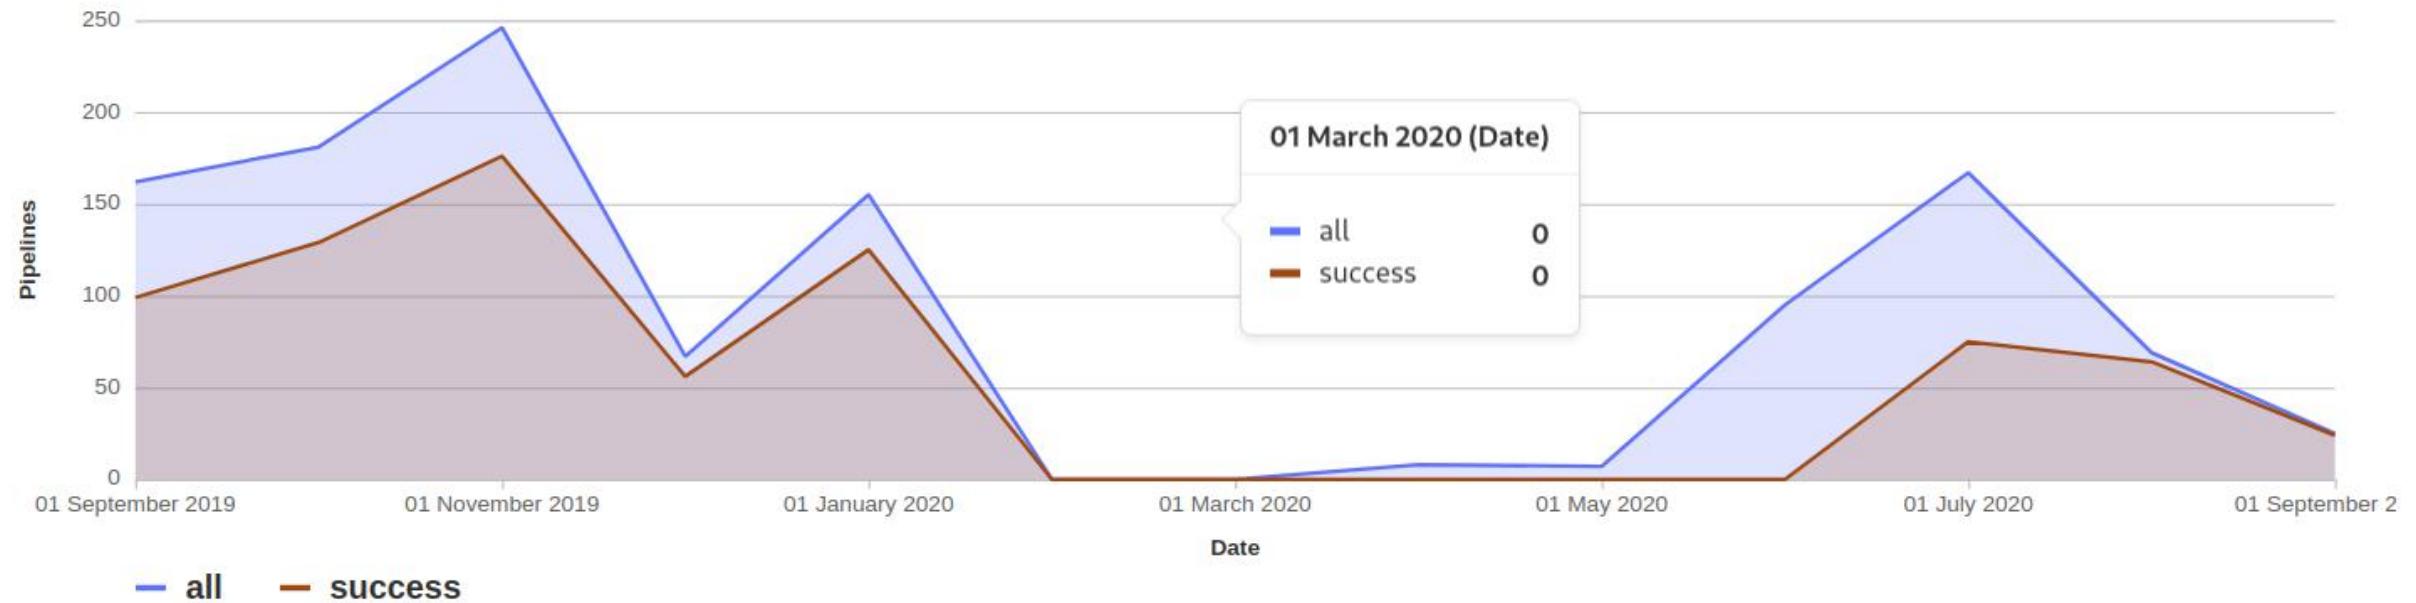
pynvme builds your own tests.

Quality: test of the test infrastructure



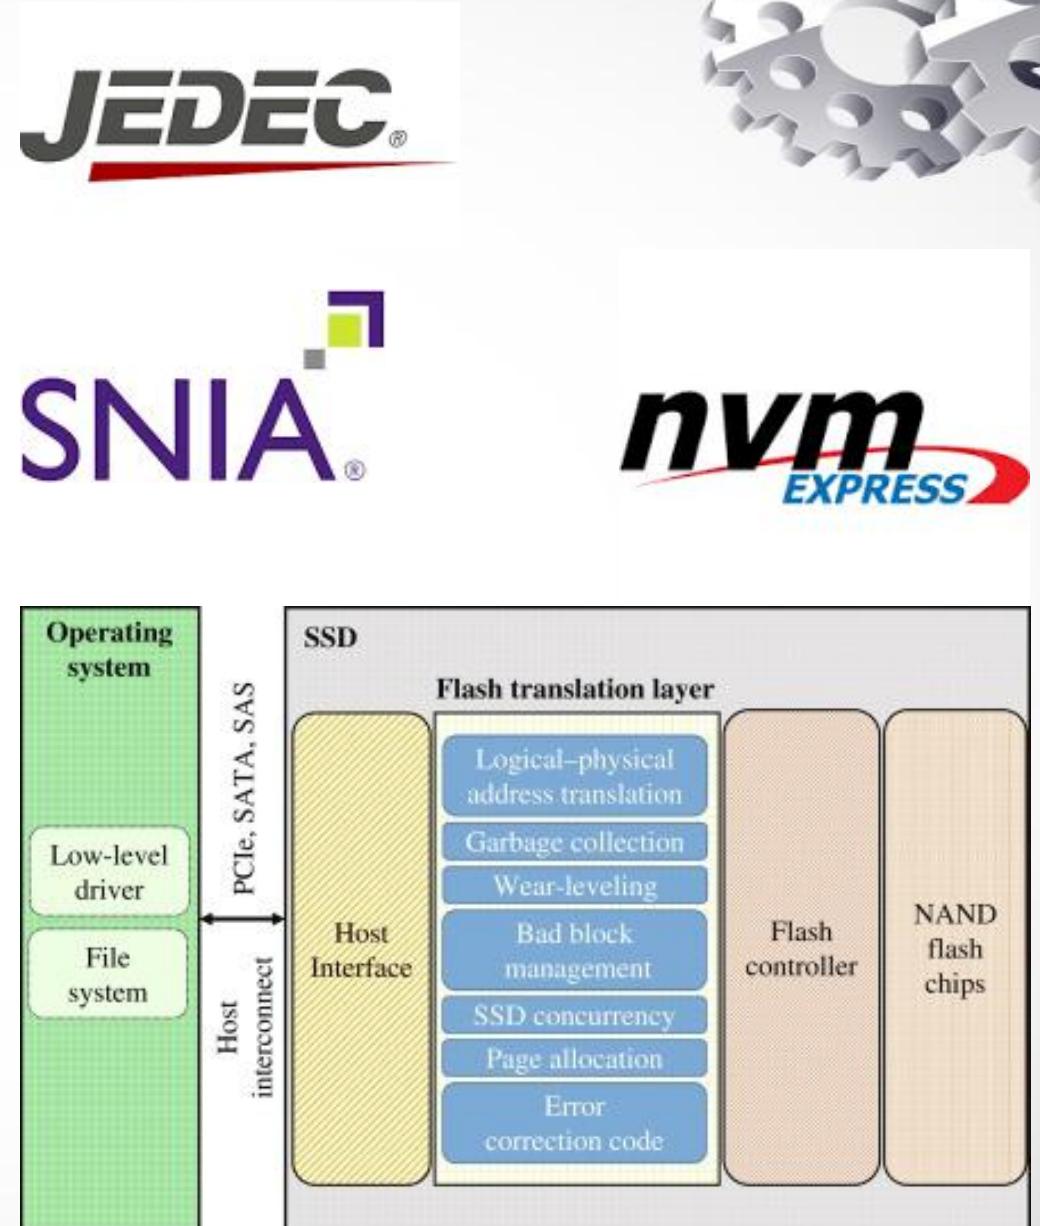
- CI in gitlab.com
- ~600 test items

Pipelines for last year



Services

```
5 def test_ioworker_jedec_workload(nvme0n1):
6     distribution = [1000]*5 + [200]*15 + [25]*80
7     iosz_distribution = {1: 4,
8                          2: 1,
9                          3: 1,
10                         4: 1,
11                         5: 1,
12                         6: 1,
13                         7: 1,
14                         8: 67,
15                         16: 10,
16                         32: 7,
17                         64: 3,
18                         128: 3}
19
20 nvme0n1.ioworker(io_size=iosz_distribution,
21                   lba_random=True,
22                   qdepth=128,
23                   distribution = distribution,
24                   read_percentage=0,
25                   ptype=0xbeef, pvalue=100,
26                   time=12*3600).start().close()
```



Conformance Test Scripts



- <https://github.com/pynvme/conformance>

- GPL-3.0

- suites

- registers
 - admin
 - nvm
 - features
 - controller
 - TCG
 - ZNS

GYTech 耕耘科技



- providing commercial services customers
- the gap between SSD development and test
 - many SSD controller and product vendors
 - SSD fast growing
- the gap between SSD vendors and customers
 - 3rd party test suites and technology supplier

Client SSD



- open source product:
 - pynvme: NVMe test driver, 2.x
 - conformance: test scripts
- commercial service to customers
 - customize pynvme
 - porting test framework
 - developing test scripts
 - training and consultation

Enterprise SSD



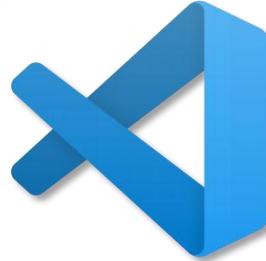
- commercial tools and services

	pynvme supported	SPDK supported	to be developed
multiple controllers	✓		
multiple namespaces	✓		
zoned namespace	✓		
namespace mgmt		✓	
reservation command		✓	
SGL		✓	
metadata		✓	
E2E data protection			✓
CMB		✓	
SRIOV			✓
latency histogram	✓		

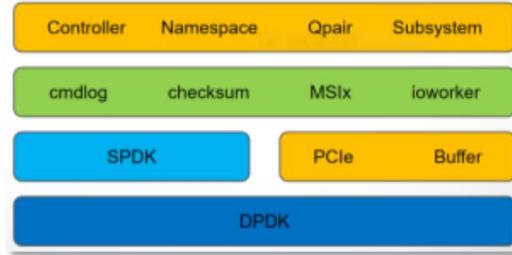
Values



Ecosystem



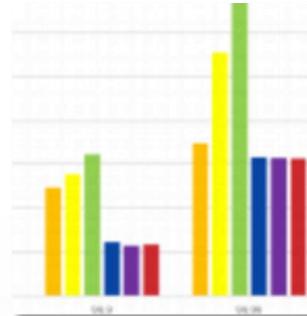
Scripts



Expendability



Hardware



Performance



Service





pynvme builds your own tests.



<https://github.com/pynvme/pynvme>

Thanks!